

CS 5814 Homework 4, Part 1: Generative Adversarial Networks

```
from google.colab import files
import zipfile
import os
# Create a zip file containing the folder you want to upload
folder_to_upload = "/path/to/your/folder"
zip_file_name = "folder_to_upload.zip"
with zipfile.ZipFile(zip_file_name, 'w', zipfile.ZIP_DEFLATED) as zipf:
    for root, dirs, files in os.walk(folder_to_upload):
        for file in files:
            zipf.write(os.path.join(root, file), os.path.relpath(os.path.join(root, file)), f
# Upload the zip file
uploaded = files.upload()
```

Choose Files cats.zip

- **cats.zip**(application/x-zip-compressed) - 49750115 bytes, last modified: 4/27/2024 - 100% done
Saving cats.zip to cats.zip

```
import zipfile
import os

# Specify the path to the zip file you uploaded
zip_file_path = "/content/cats.zip"

# Specify the directory where you want to extract and save the contents
extracted_folder_path = "/content/cats"

# Create a directory to extract the contents if it doesn't exist
if not os.path.exists(extracted_folder_path):
    os.makedirs(extracted_folder_path)

# Extract and save the contents of the zip file
with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
    zip_ref.extractall(extracted_folder_path)

print("Files extracted and saved successfully to:", extracted_folder_path)

Files extracted and saved successfully to: /content/cats
```

```
import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder
import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

# %pip install torchvision
```

WARNING: The directory '/home/jovyan/.cache/pip' or its parent directory is not owned or Defaulting to user installation because normal site-packages is not writeable
Collecting torchvision

 Downloading torchvision-0.17.2-cp38-cp38-manylinux1_x86_64.whl (6.9 MB)
 |██████████| 6.9 MB 3.3 MB/s eta 0:00:01

Requirement already satisfied: numpy in /opt/conda/lib/python3.8/site-packages (from tor
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /opt/conda/lib/python3.8/site-pa
Requirement already satisfied: torch==2.2.2 in /home/stanleygabriel97/.local/lib/python3.
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.1.105; platform_system == "Lir
Requirement already satisfied: nvidia-curand-cu12==10.3.2.106; platform_system == "Linux"
Requirement already satisfied: nvidia-nccl-cu12==2.19.3; platform_system == "Linux" and
Requirement already satisfied: triton==2.2.0; platform_system == "Linux" and platform_ma
Requirement already satisfied: nvidia-cusparse-cu12==12.1.0.106; platform_system == "Lir
Requirement already satisfied: sympy in /opt/conda/lib/python3.8/site-packages (from tor
Requirement already satisfied: networkx in /opt/conda/lib/python3.8/site-packages (from
Requirement already satisfied: nvidia-cufft-cu12==11.0.2.54; platform_system == "Linux"
Requirement already satisfied: typing-extensions>=4.8.0 in /home/stanleygabriel97/.local
Requirement already satisfied: nvidia-nvtx-cu12==12.1.105; platform_system == "Linux" ar
Requirement already satisfied: jinja2 in /opt/conda/lib/python3.8/site-packages (from tc
Requirement already satisfied: filelock in /home/stanleygabriel97/.local/lib/python3.8/s
Requirement already satisfied: nvidia-cublas-cu12==12.1.3.1; platform_system == "Linux"
Requirement already satisfied: nvidia-cusolver-cu12==11.4.5.107; platform_system == "Lir
Requirement already satisfied: fsspec in /opt/conda/lib/python3.8/site-packages (from tc
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.1.105; platform_system == "L
Requirement already satisfied: nvidia-cudnn-cu12==8.9.2.26; platform_system == "Linux" a
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.1.105; platform_system == "Lir
Requirement already satisfied: nvidia-nvjitlink-cu12 in /home/stanleygabriel97/.local/li
Requirement already satisfied: mpmath>=0.19 in /opt/conda/lib/python3.8/site-packages (f
Requirement already satisfied: decorator>=4.3.0 in /opt/conda/lib/python3.8/site-package
Requirement already satisfied: MarkupSafe>=0.23 in /opt/conda/lib/python3.8/site-package
Installing collected packages: torchvision
Successfully installed torchvision-0.17.2
Note: you may need to restart the kernel to use updated packages.

```
from google.colab import drive
drive.mount("/content/drive")

Mounted at /content/drive

# import os
# cwd = os.getcwd()

# cwd
'/home/stanleygabriel97/cs5814/hw4/part1'

# import os
# datadir = "./part1/" # path to the homework
# if not os.path.exists(cwd):
#   !ln -s "$datadir" # path to the homework
# os.chdir(cwd)
# !pwd

/home/stanleygabriel97/cs5814/hw4/part1

import zipfile
import os

# Specify the path to the zip file you uploaded
zip_file_path = "/content/cats.zip"

# Specify the directory where you want to extract and save the contents
extracted_folder_path = "/content/cats"

# Create a directory to extract the contents if it doesn't exist
if not os.path.exists(extracted_folder_path):
    os.makedirs(extracted_folder_path)

# Extract and save the contents of the zip file
with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
    zip_ref.extractall(extracted_folder_path)

print("Files extracted and saved successfully to:", extracted_folder_path)

Files extracted and saved successfully to: /content/cats

import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder
```

```
import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

The autoreload extension is already loaded. To reload it, use:
%reload_ext autoreload

uploaded = files.upload()
```

Choose Files 4 files

- **losses.py**(text/x-python) - 3747 bytes, last modified: 5/1/2024 - 100% done
 - **models.py**(text/x-python) - 2021 bytes, last modified: 5/1/2024 - 100% done
 - **train.py**(text/x-python) - 4489 bytes, last modified: 5/1/2024 - 100% done
 - **utils.py**(text/x-python) - 1609 bytes, last modified: 5/1/2024 - 100% done
- Saving losses.py to losses.py
Saving models.py to models.py
Saving train.py to train.py
Saving utils.py to utils.py

```
from train import train

device = torch.device("cuda:0" if torch.cuda.is_available() else "gpu")

device
device(type='cuda', index=0)
```

▼ GAN loss functions

You'll need to implement two different loss functions. The first is the loss from the [original GAN paper](#). The next is the loss from [LS-GAN](#).

▼ GAN loss

TODO: You need to implement the `discriminator_loss` and `generator_loss` functions in `gan/losses.py`.

The generator loss is given by:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

Note that these equations are negated because we will be *minimizing* these losses.

HINTS: Use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss.

We will be averaging over the elements of the minibatch instead of computing the expectation of $\log D(G(z))$, $\log D(x)$ and $\log(1 - D(G(z)))$.

```
from losses import discriminator_loss, generator_loss
```

✓ Least Squares GAN loss

TODO: You need to implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

[Least Squares GAN](#) is an alternative to the original GAN loss function. For this part, all you need to do change the loss function and retrain the model.

Specifically, you'll implement equation (9) in the paper:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

```
from losses import ls_discriminator_loss, ls_generator_loss
```

✓ GAN model architecture

TODO: Next, you'll need to implement the Discriminator and Generator networks in `gan/models.py`.

We'll be using the architecture from [DCGAN](#):

Discriminator:

- convolutional layer with `in_channels=3, out_channels=128, kernel=4, stride=2`

- convolutional layer with in_channels=128, out_channels=256, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=256, out_channels=512, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=512, out_channels=1024, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=1024, out_channels=1, kernel=4, stride=1

Use padding = 1 (not 0) for all the convolutional layers.

You can either use LeakyReLu throughout the discriminator (and a negative slope value of 0.2) or just use relu.

The discriminator will output a single score for each sample. The output of your discriminator should be a single value score corresponding to each input sample.

Generator:

Note: Here, you'll need to use transposed convolution (sometimes known as fractionally-strided convolution). This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1
- batch norm
- transpose convolution with in_channels=1024, out_channels=512, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=512, out_channels=256, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=256, out_channels=128, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=128, out_channels=3, kernel=4, stride=2

The generator network will need to map the output values between -1 and 1 using a `tanh` nonlinearity. The output should be of size 64x64 with 3 channels for each sample (equal dimensions to the images from the dataset).

```
from models import Discriminator, Generator
```

▼ Data loading

Our dataset is pretty small, so in order to prevent overfitting, we need to perform data augmentation.

TODO: You'll need to implement some data augmentation by adding new transforms to the cell below.

The easiest you can do is just use the RandomCrop and ColorJitter, but feel free to play around with other augmentations as well to see how it affects the performance of your model. See <https://pytorch.org/vision/stable/transforms.html>.

```
batch_size = 32
imsize = 64
cat_root = '/content/cats'

# Define your transformations
transform_augment = transforms.Compose([
    transforms.Resize((imsize, imszie)), # Ensures all images are the same size
    transforms.RandomCrop(imszie, padding=4), # Pads and then crops
    transforms.ColorJitter(brightness=0.1, contrast=0.1, saturation=0.1, hue=0.1),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomRotation(15),
    transforms.ToTensor(), # Converts to tensor
])

# Create your ImageFolder dataset
cat_train = ImageFolder(root=cat_root, transform=transform_augment)

# Create DataLoader
cat_loader_train = DataLoader(cat_train, batch_size=batch_size, shuffle=True, drop_last=True)
```

▼ Visualize dataset

```
from gan.utils import show_images

for batch in cat_loader_train:
    imgs = batch[0].numpy().squeeze()
    show_images(imgs, color=True)
    break
```



▼ Training

TODO: You need to write the training loop in `gan/train.py`.

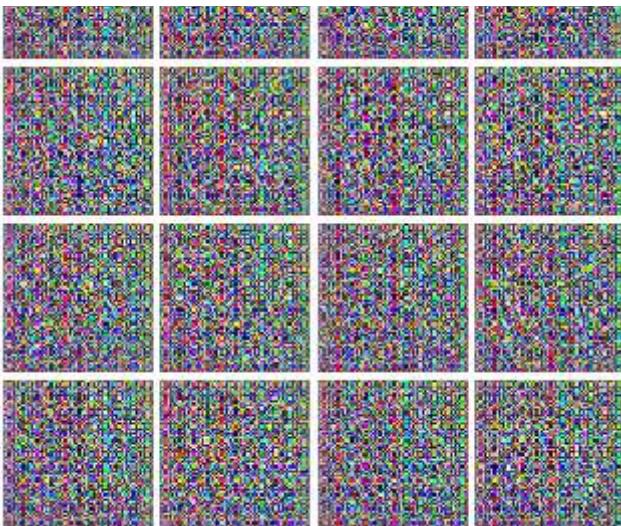
```
NOISE_DIM = 100  
NUM_EPOCHS = 50  
learning_rate = 0.001
```

▼ Train GAN

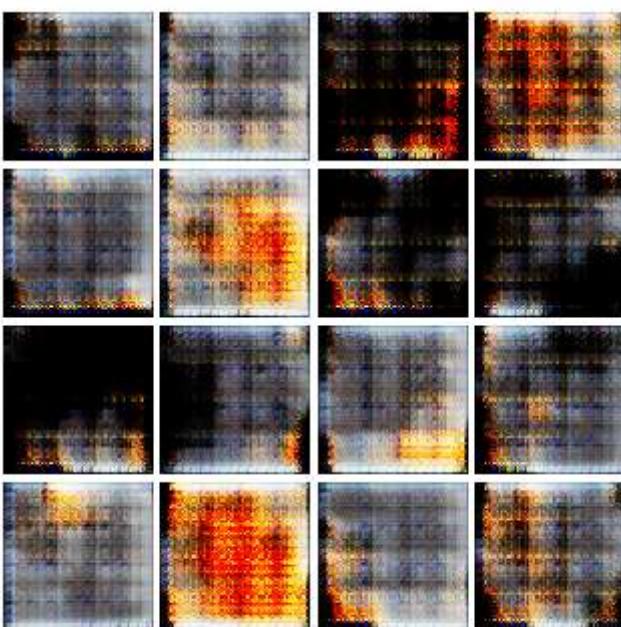
```
D = Discriminator().to(device)  
G = Generator(noise_dim=NOISE_DIM).to(device)
```

```
D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))  
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))
```

```
# original gan
train(D, G, D_optimizer, G_optimizer, discriminator_loss,
      generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

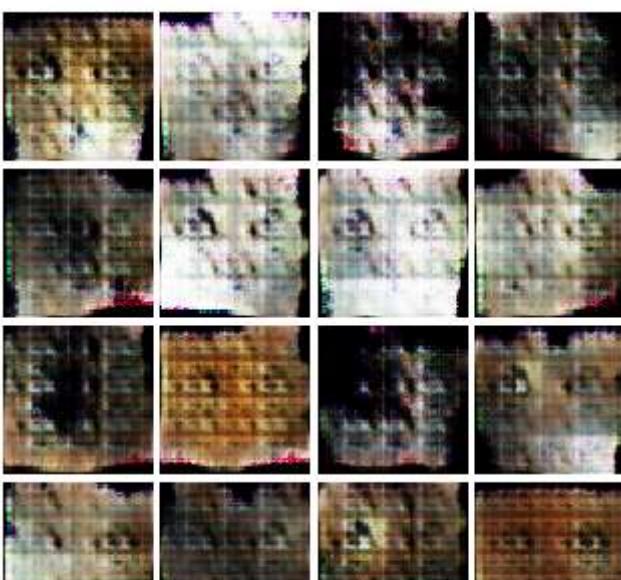


Iter: 250, D: 1.032, G:4.5



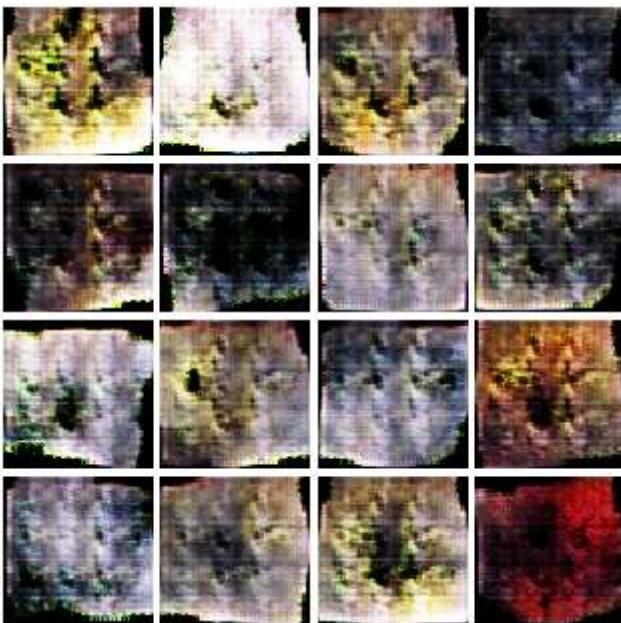
EPOCH: 2

Iter: 500, D: 0.8907, G:1.909



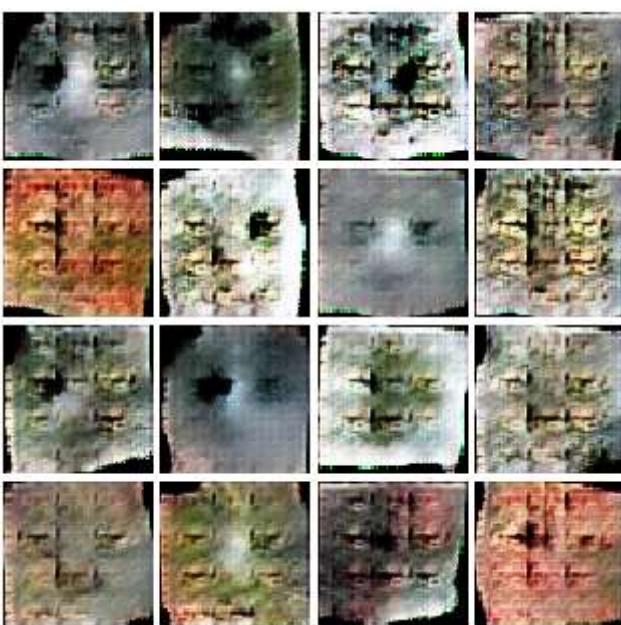


Iter: 750, D: 0.9202, G:3.247



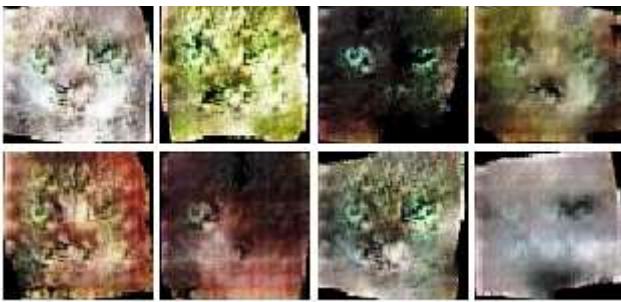
EPOCH: 3

Iter: 1000, D: 1.417, G:1.554



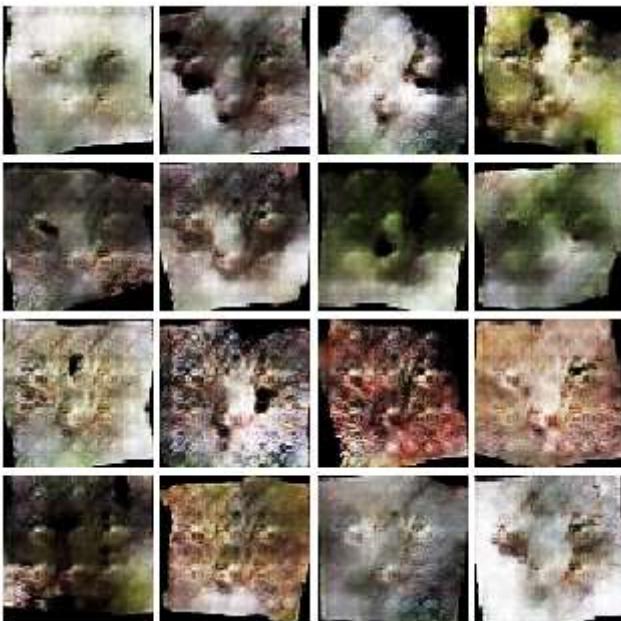
Iter: 1250, D: 1.204, G:3.537





EPOCH: 4

Iter: 1500, D: 1.536, G:2.177



Iter: 1750, D: 1.267, G:4.24



EPOCH: 5

Iter: 2000, D: 1.007, G:1.771



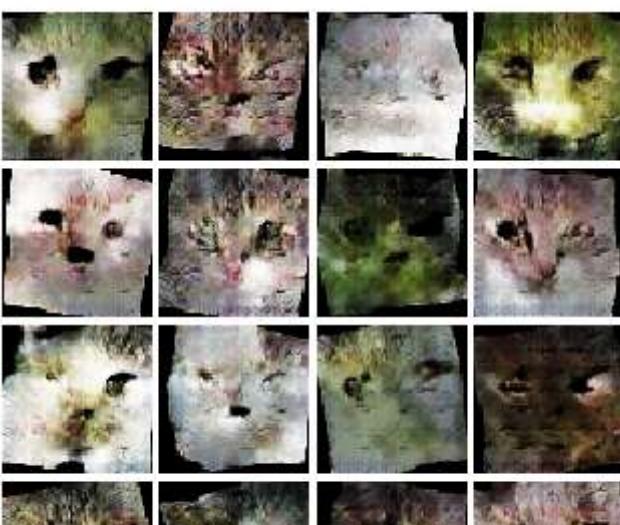


Iter: 2250, D: 1.022, G:2.23



EPOCH: 6

Iter: 2500, D: 1.305, G:3.021





Iter: 2750, D: 0.9092, G:2.5



EPOCH: 7

Iter: 3000, D: 0.6024, G:3.235



Iter: 3250, D: 0.6535, G:3.135





EPOCH: 8

Iter: 3500, D: 0.6437, G:4.008



Iter: 3750, D: 1.176, G:5.605



EPOCH: 9

Iter: 4000, D: 0.6453, G:4.451



Iter: 4250, D: 0.7271, G: 4.437



EPOCH: 10

Iter: 4500, D: 0.5916, G: 3.872





Iter: 4750, D: 0.5339, G: 3.419



EPOCH: 11

Iter: 5000, D: 0.5941, G: 3.317



Iter: 5250, D: 1.549, G: 2.761





EPOCH: 12

Iter: 5500, D: 0.3588, G:4.996



Iter: 5750, D: 0.6169, G:6.276



Epoch: 12

EPOCH: 15

Iter: 6000, D: 0.6193, G:5.159



Iter: 6250, D: 0.681, G:10.27



EPOCH: 14

Iter: 6500, D: 0.3038, G:5.288





Iter: 6750, D: 0.1952, G: 3.823



EPOCH: 15

Iter: 7000, D: 0.4229, G: 6.833



Iter: 7250, D: 0.1874, G: 5.001





EPOCH: 16

Iter: 7500, D: 0.1565, G: 6.137



Iter: 7750, D: 0.9893, G: 2.385



EPOCH: 17

Iter: 8000, D: 0.3265, G:6.547



Iter: 8250, D: 0.4314, G:6.103



EPOCH: 18

Iter: 8500, D: 0.06536, G:5.327





Iter: 8750, D: 0.2058, G:5.863



EPOCH: 19

Iter: 9000, D: 0.2069, G:8.763



Iter: 9250, D: 0.8755, G:11.01





EPOCH: 20

Iter: 9500, D: 0.3802, G: 8.224



Iter: 9750, D: 0.03611, G: 6.363





EPOCH: 21

Iter: 10000, D: 0.1062, G:7.084



Iter: 10250, D: 0.0477, G:6.773



EPOCH: 22

Iter: 10500, D: 0.2313, G:5.054





Iter: 10750, D: 0.07899, G:5.622



EPOCH: 23

Iter: 11000, D: 0.1112, G:6.658



Iter: 11250, D: 1.728, G:21.38





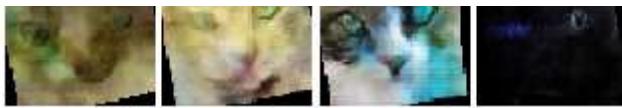
EPOCH: 24

Iter: 11500, D: 0.03007, G:6.568



Iter: 11750, D: 0.2212, G:5.852





EPOCH: 25

Iter: 12000, D: 0.09737, G:7.974



Iter: 12250, D: 0.06177, G:4.739



EPOCH: 26

Iter: 12500, D: 0.08106, G:6.075





Iter: 12750, D: 0.2877, G:6.02



EPOCH: 27

Iter: 13000, D: 0.04758, G:6.504



Iter: 13250, D: 0.08282, G:7.503



EPOCH: 28

Iter: 13500, D: 0.3406, G: 9.794



Iter: 13750, D: 0.02785, G: 7.07





EPOCH: 29

Iter: 14000, D: 0.02731, G:5.526



Iter: 14250, D: 0.3407, G:4.559



EPOCH: 30

Iter: 14500, D: 0.1133, G:4.504





Iter: 14750, D: 0.5945, G: 3.644



EPOCH: 31

Iter: 15000, D: 0.246, G: 7.006



Iter: 15250, D: 0.06481, G:8.699



EPOCH: 32

Iter: 15500, D: 0.09275, G:5.653



EPOCH: 33

Iter: 15750, D: 0.1036, G:10.19





Iter: 16000, D: 0.4116, G: 9.544



EPOCH: 34

Iter: 16250, D: 0.08335, G: 6.745



Iter: 16500, D: 1.241, G: 3.693





EPOCH: 35

Iter: 16750, D: 0.03554, G:12.18



Iter: 17000, D: 0.6134, G:8.738



EPOCH: 36
Iter: 17250, D: 0.0218, G:7.417



Iter: 17500, D: 1.015, G:15.61



EPOCH: 37
Iter: 17750, D: 0.008697, G:11.27





Iter: 18000, D: 0.06906, G:6.486



EPOCH: 38

Iter: 18250, D: 0.01343, G:7.557



Iter: 18500, D: 0.2975, G:10.04





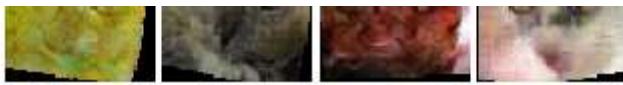
EPOCH: 39

Iter: 18750, D: 0.03125, G:7.357



Iter: 19000, D: 0.2797, G:6.413





EPOCH: 40

Iter: 19250, D: 0.04077, G:10.37



Iter: 19500, D: 0.04365, G:8.713



EPOCH: 41

Iter: 19750, D: 0.0787, G:6.169





Iter: 20000, D: 0.2463, G: 4.271



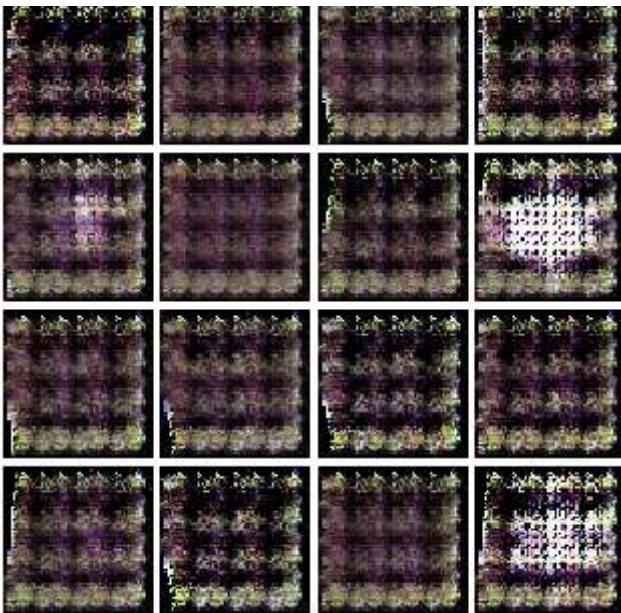
EPOCH: 42

Iter: 20250, D: 0.03496, G: 22.24

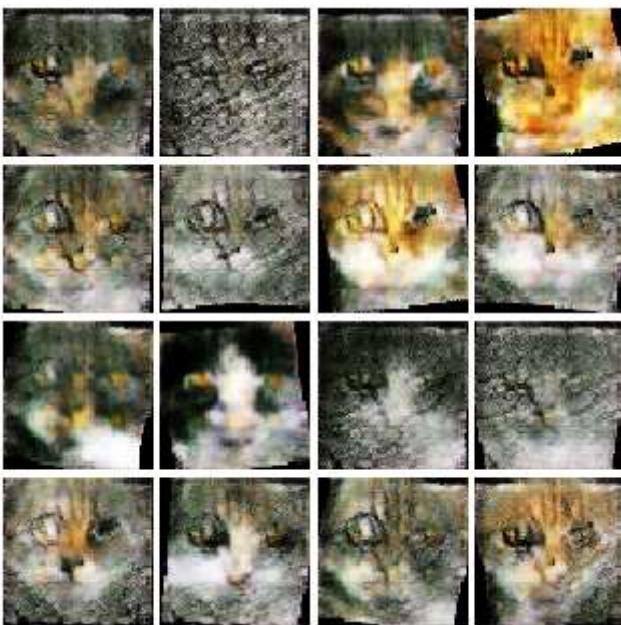


Iter: 20500, D: 6.075e-05, G: 28.46





EPOCH: 43
Iter: 20750, D: 0.03668, G:12.67



Iter: 21000, D: 0.08736, G:5.776



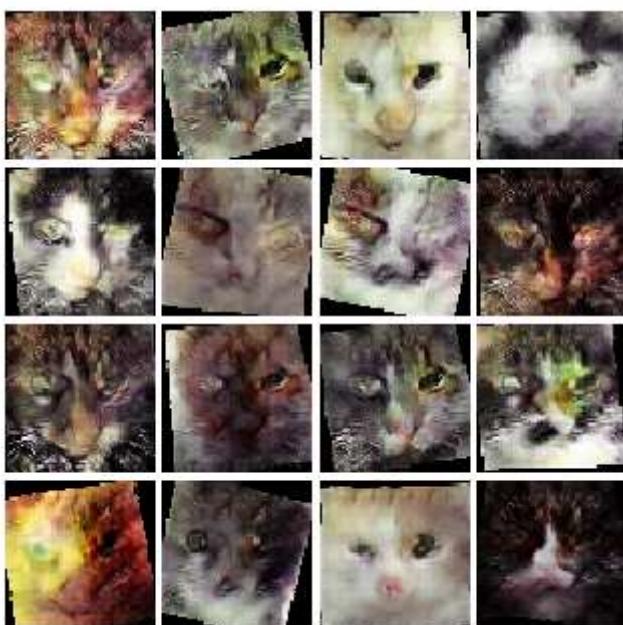


EPOCH: 44

Iter: 21250, D: 0.3015, G:5.206



Iter: 21500, D: 0.00915, G:7.408



EPOCH: 45

Iter: 21750, D: 0.0346, G:7.252





Iter: 22000, D: 0.2835, G: 8.013



EPOCH: 46

Iter: 22250, D: 0.09855, G: 11.97

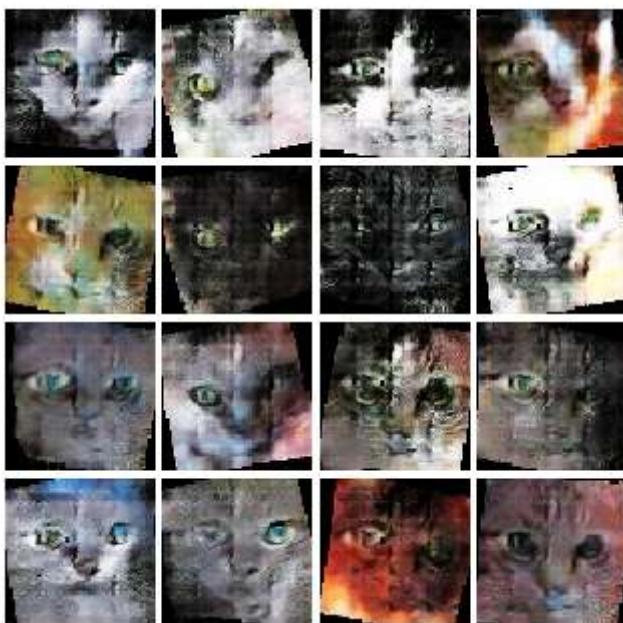


Iter: 22500, D: 0.04589, G:9.738



EPOCH: 47

Iter: 22750, D: 0.2167, G:10.64



Iter: 23000, D: 0.5244, G:16.5





EPOCH: 48

Iter: 23250, D: 0.2772, G: 5.852



Iter: 23500, D: 0.0253, G: 9.613



EPOCH: 49

Iter: 23750, D: 0.0905, G: 4.809





Iter: 24000, D: 0.8134, G:18.92

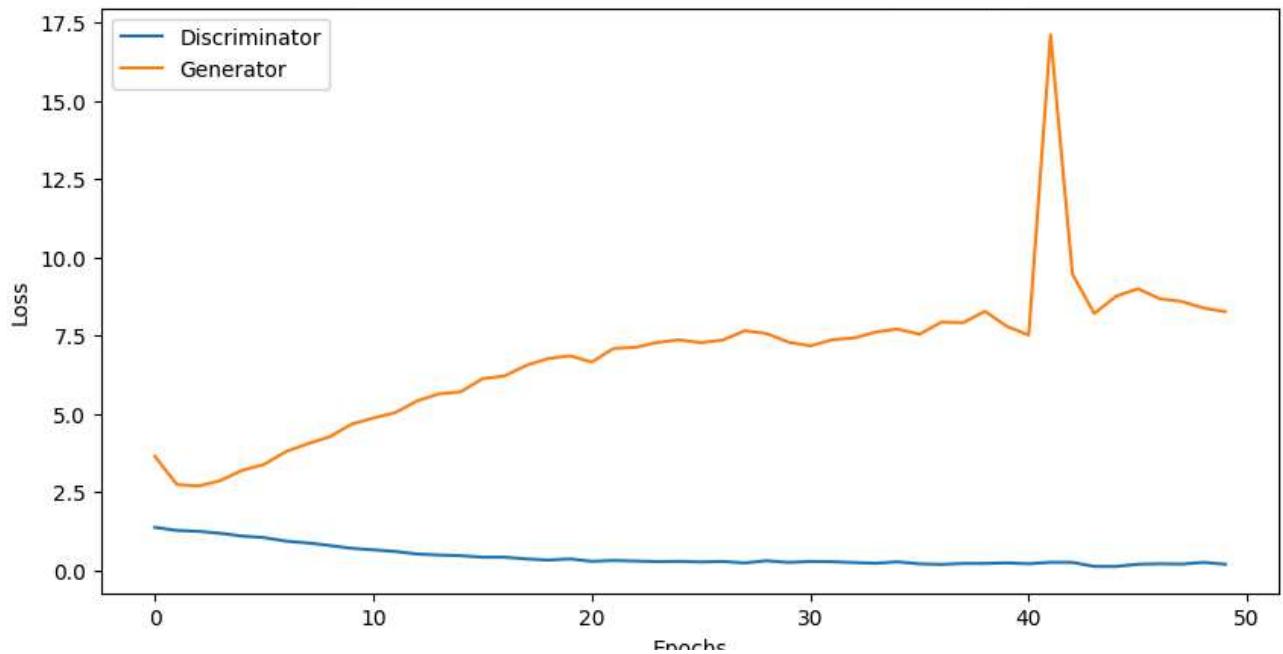


EPOCH: 50

Iter: 24250, D: 0.00371, G:11.46



Iter: 24500, D: 0.4889, G:15.0



▼ Train LS-GAN

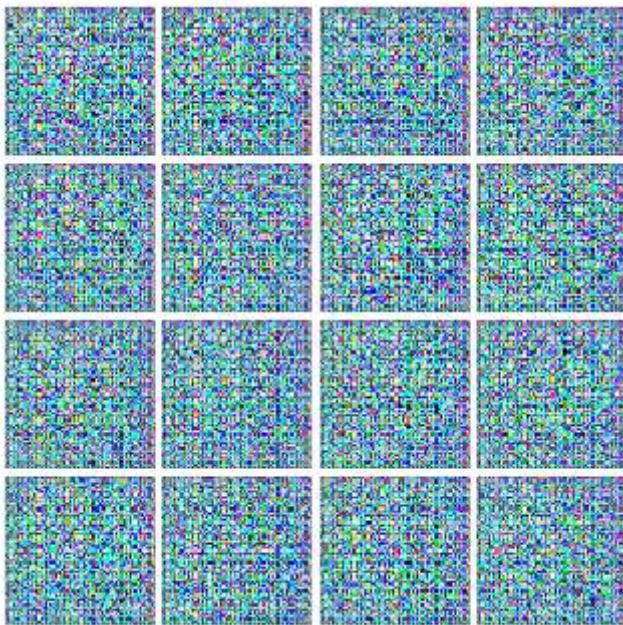
```
D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)
```

```
D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

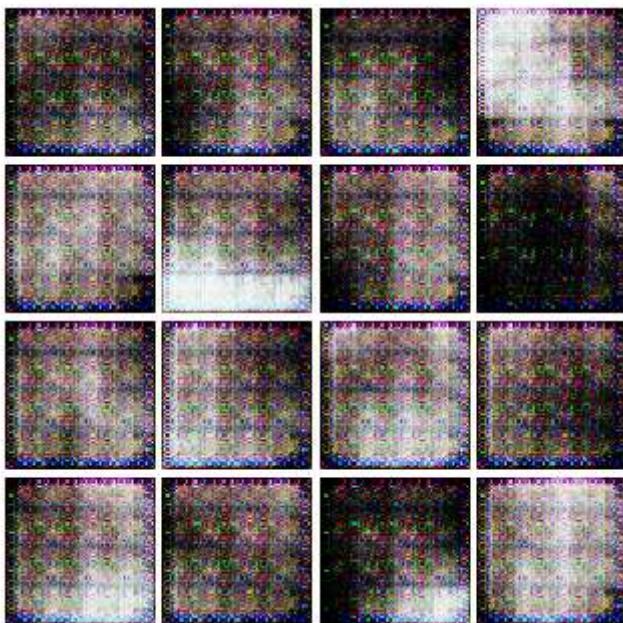
# ls-gan
train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
      ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

EPOCH: 1

Iter: 0, D: 0.7362, G:198.9

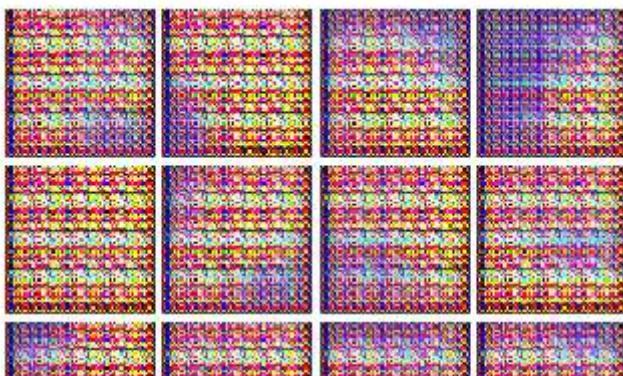


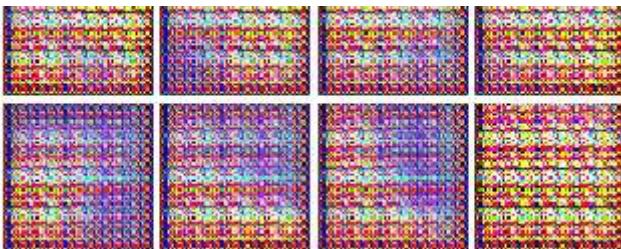
Iter: 250, D: 0.4907, G:0.5792



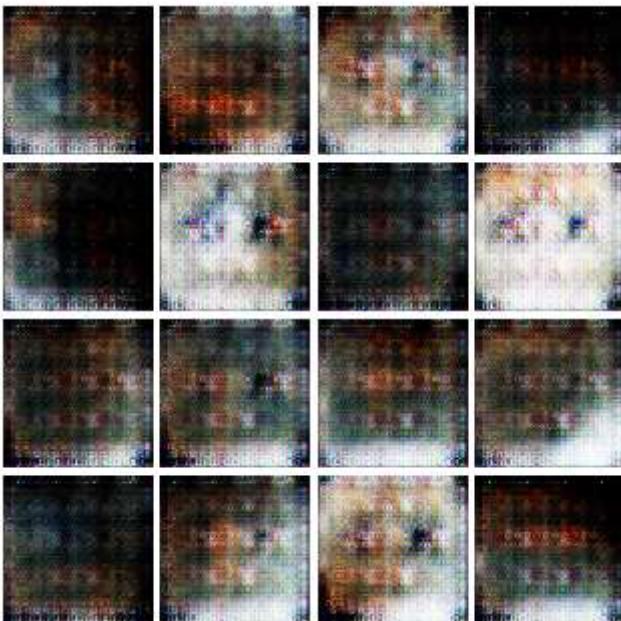
EPOCH: 2

Iter: 500, D: 0.1546, G:0.2421





Iter: 750, D: 0.2104, G:0.5662



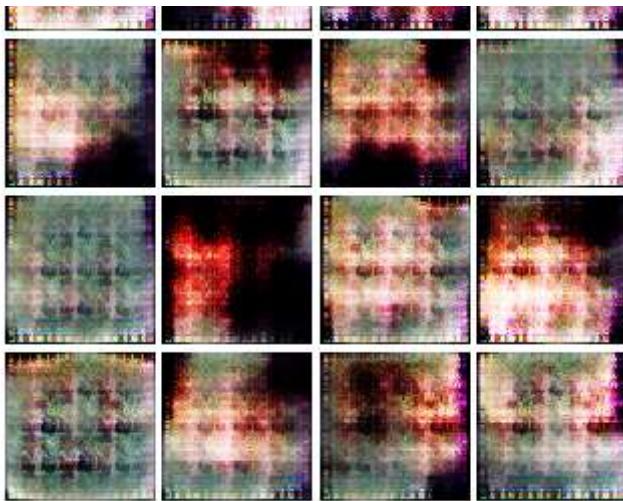
EPOCH: 3

Iter: 1000, D: 0.3954, G:0.9194



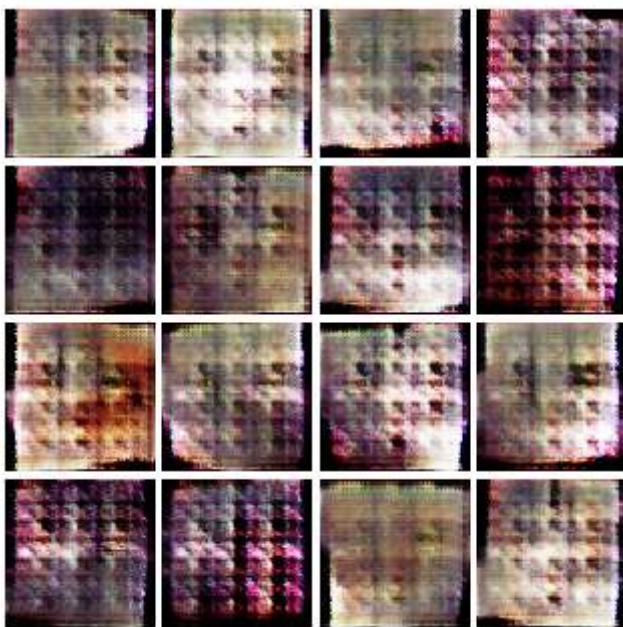
Iter: 1250, D: 0.3537, G:0.5337



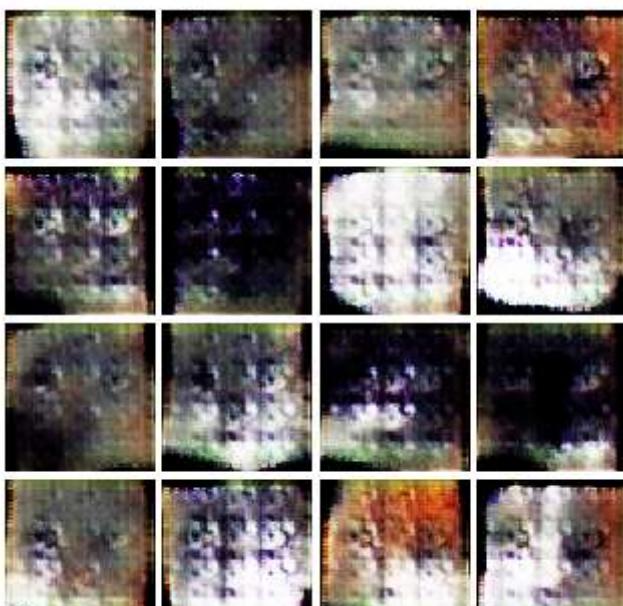


EPOCH: 4

Iter: 1500, D: 0.1804, G:0.3405



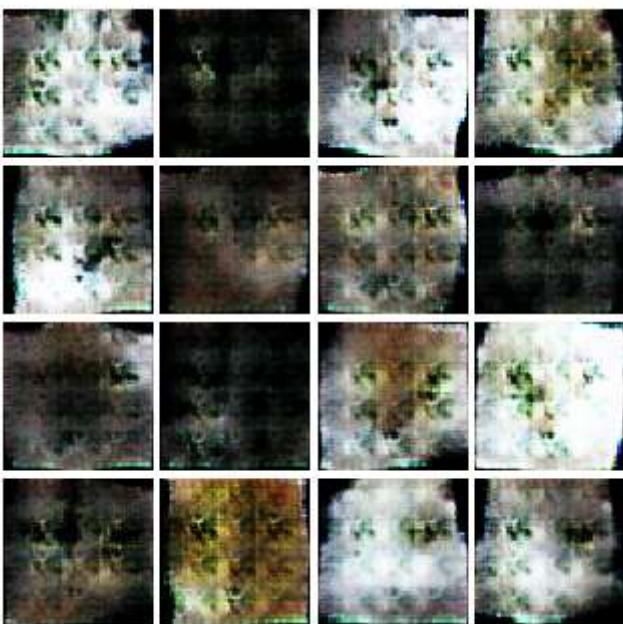
Iter: 1750, D: 0.1479, G:0.5585





EPOCH: 5

Iter: 2000, D: 0.163, G:0.5274



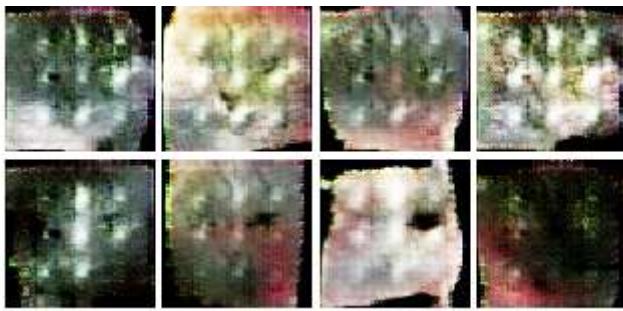
Iter: 2250, D: 0.5651, G:0.9845



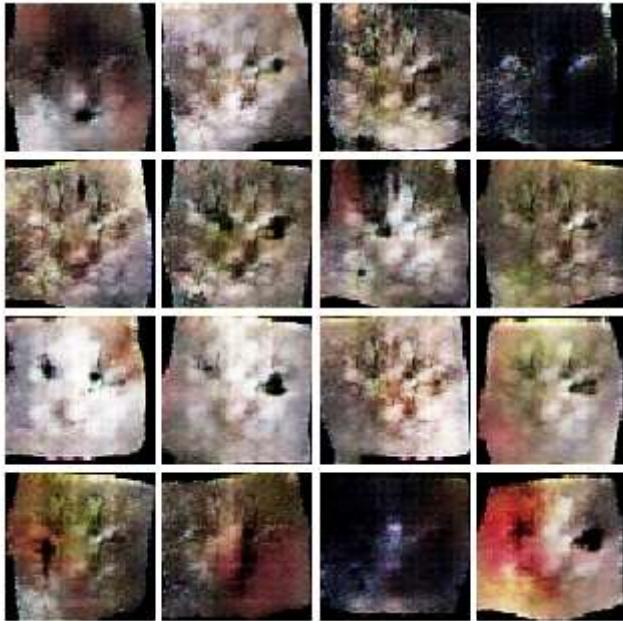
EPOCH: 6

Iter: 2500, D: 0.2795, G:0.1416



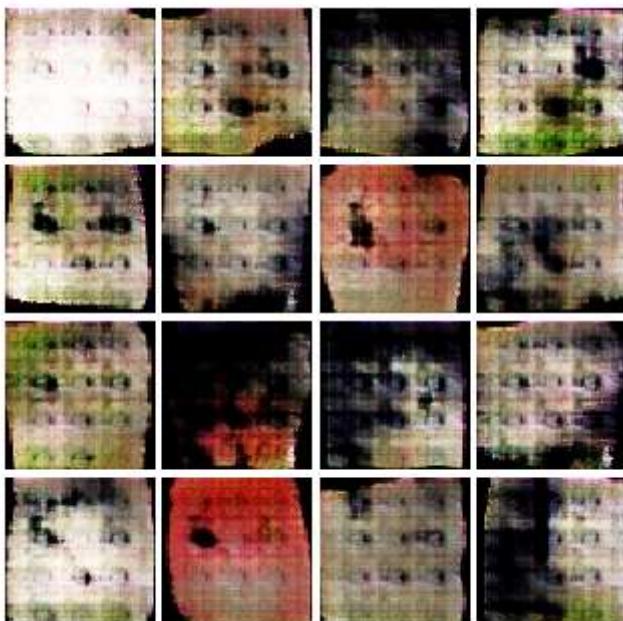


Iter: 2750, D: 0.2633, G:0.1754



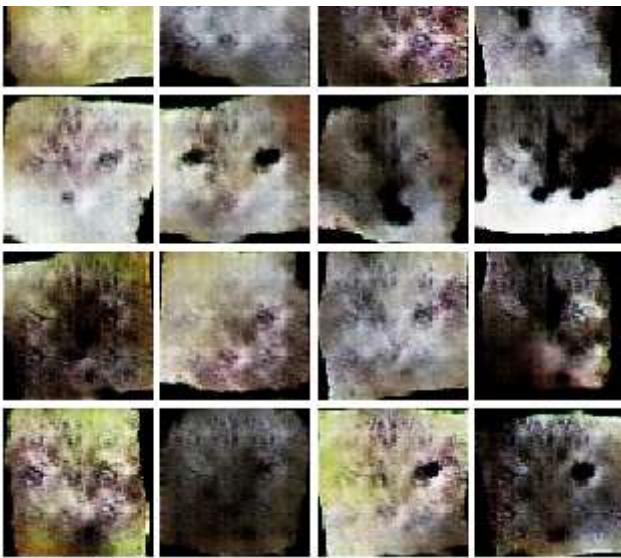
EPOCH: 7

Iter: 3000, D: 0.2583, G:0.1023



Iter: 3250, D: 0.2431, G:0.2854



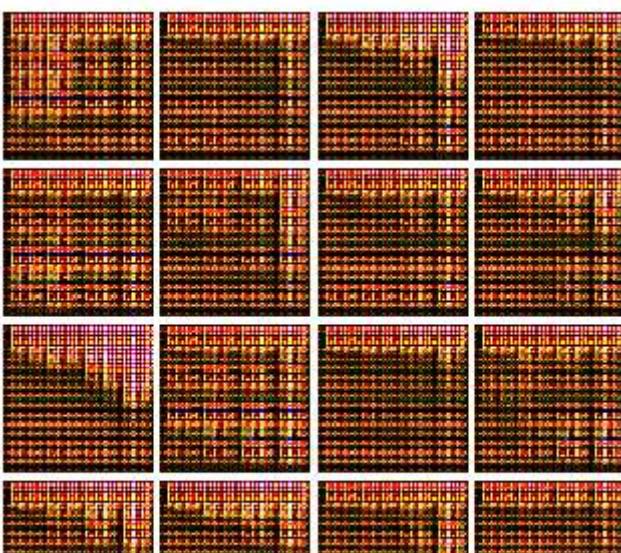


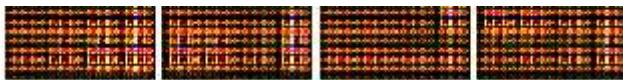
EPOCH: 8

Iter: 3500, D: 0.3165, G:0.5218



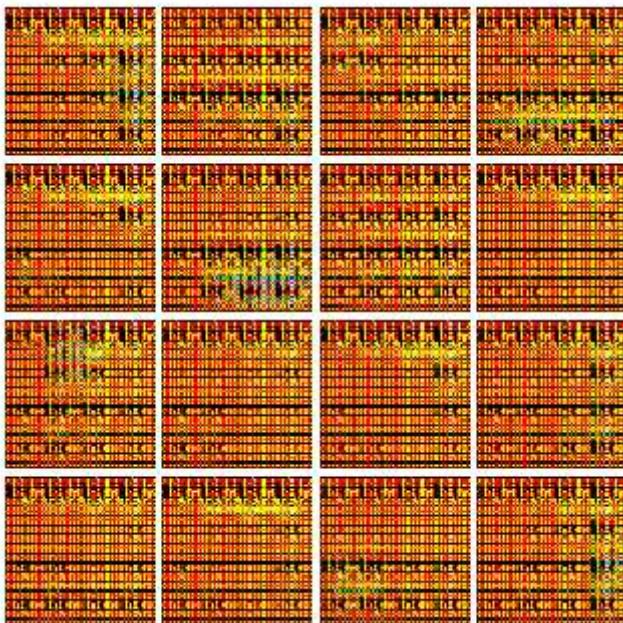
Iter: 3750, D: 0.2438, G:0.5267



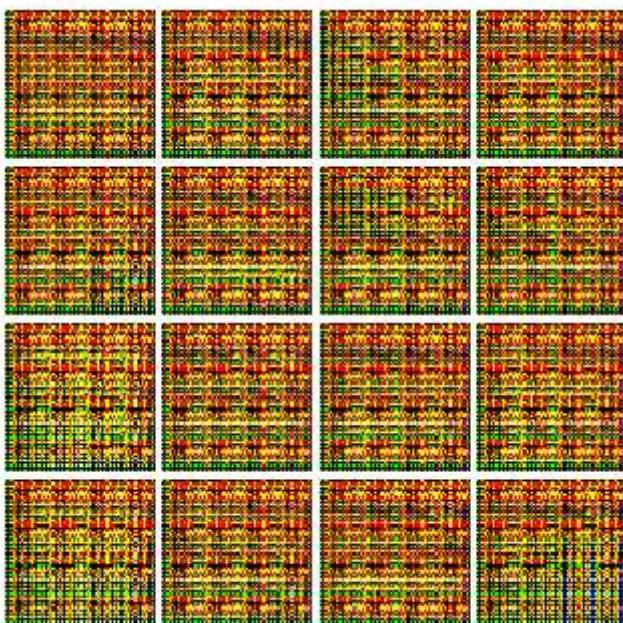


EPOCH: 9

Iter: 4000, D: 0.02441, G:0.5369

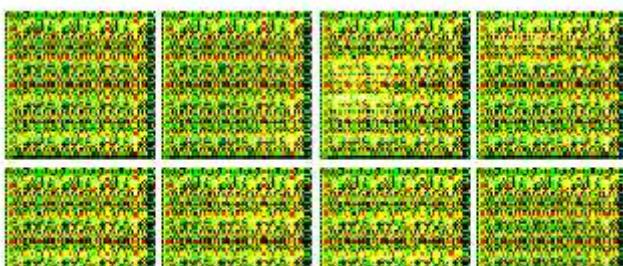


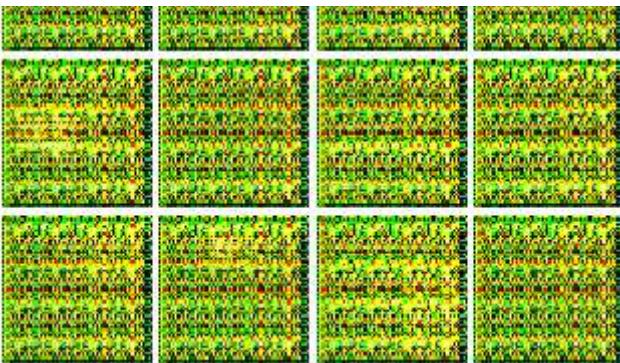
Iter: 4250, D: 0.02449, G:0.6189



EPOCH: 10

Iter: 4500, D: 0.1512, G:0.3595





Iter: 4750, D: 0.01687, G:0.5119



EPOCH: 11

Iter: 5000, D: 0.02509, G:0.4689

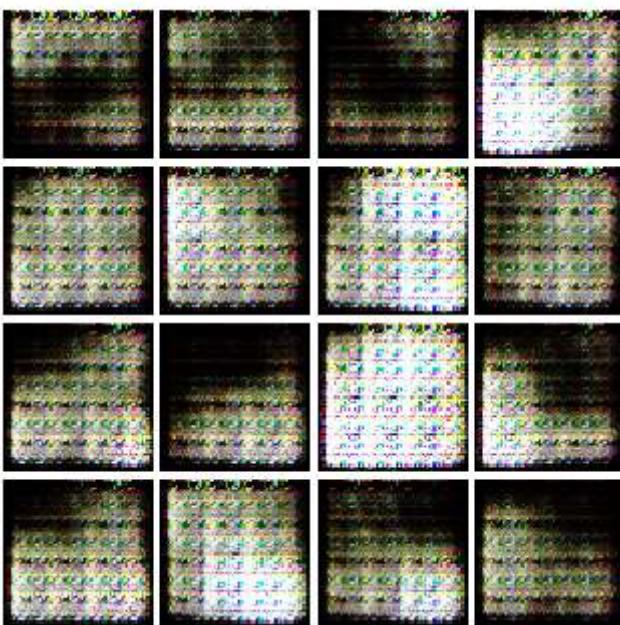


Iter: 5250, D: 0.05062, G:0.5853

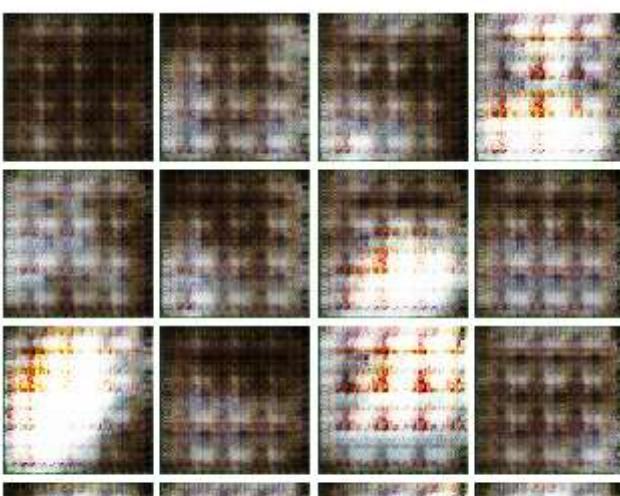


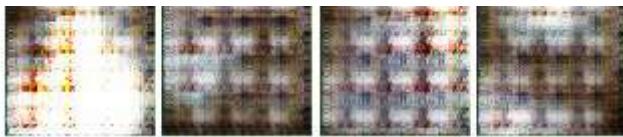
EPOCH: 12

Iter: 5500, D: 0.1076, G:0.3877



Iter: 5750, D: 0.1561, G:0.3191





EPOCH: 13

Iter: 6000, D: 0.1195, G:0.4414



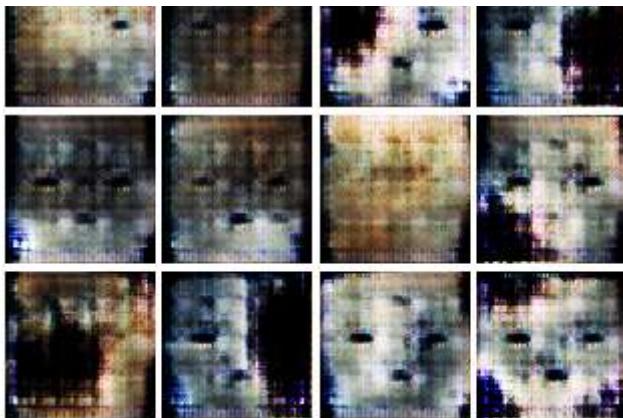
Iter: 6250, D: 0.1736, G:0.257



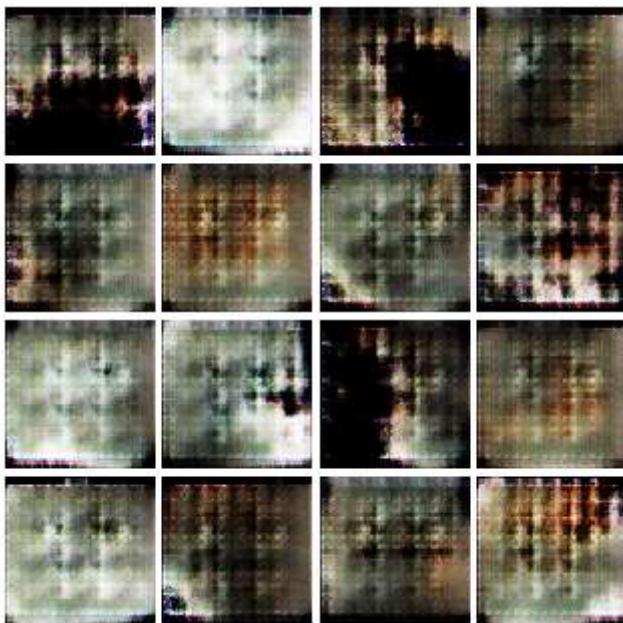
EPOCH: 14

Iter: 6500, D: 0.2123, G:0.279





Iter: 6750, D: 0.2244, G:0.3181



EPOCH: 15

Iter: 7000, D: 0.3048, G:0.5716



Iter: 7250, D: 0.2361, G:0.2867

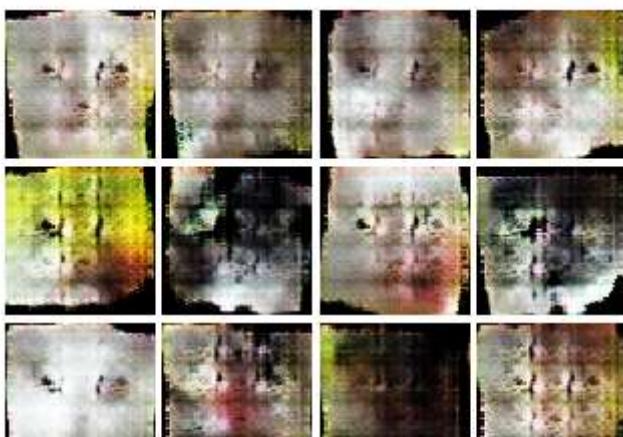


EPOCH: 16

Iter: 7500, D: 0.195, G:0.4263



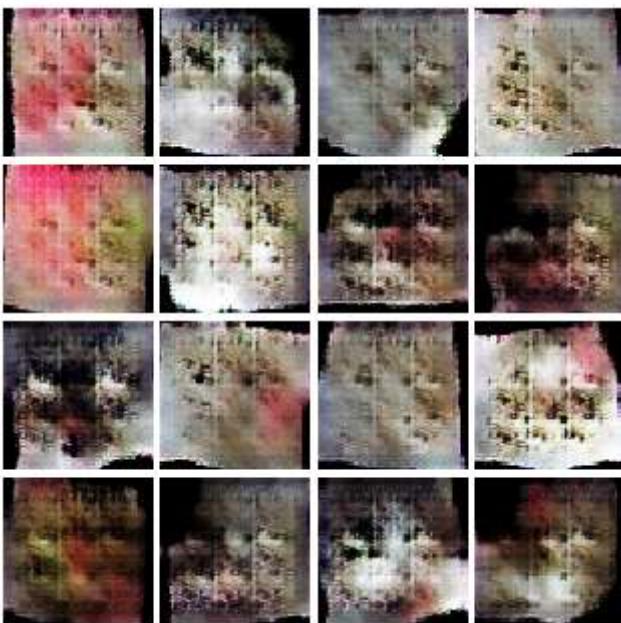
Iter: 7750, D: 0.1387, G:0.5116



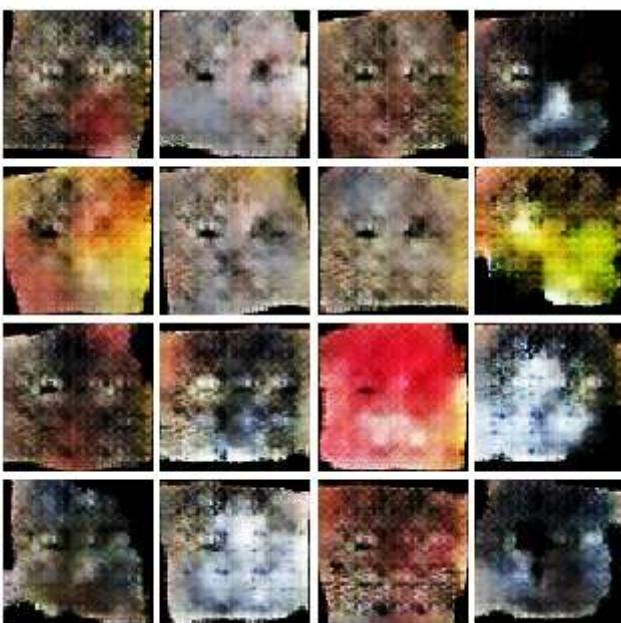


EPOCH: 17

Iter: 8000, D: 0.3599, G: 0.2998



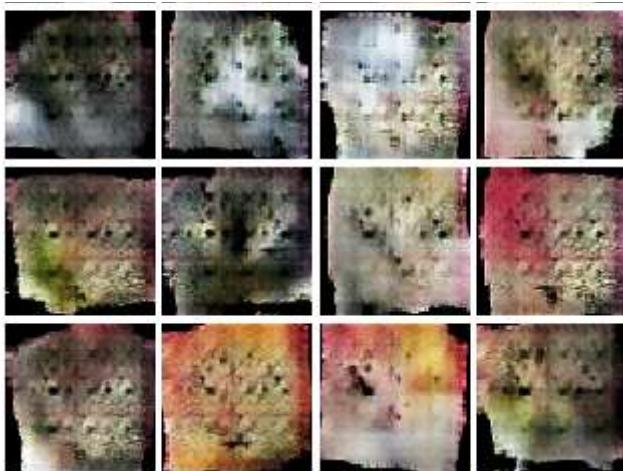
Iter: 8250, D: 1.251, G: 0.368



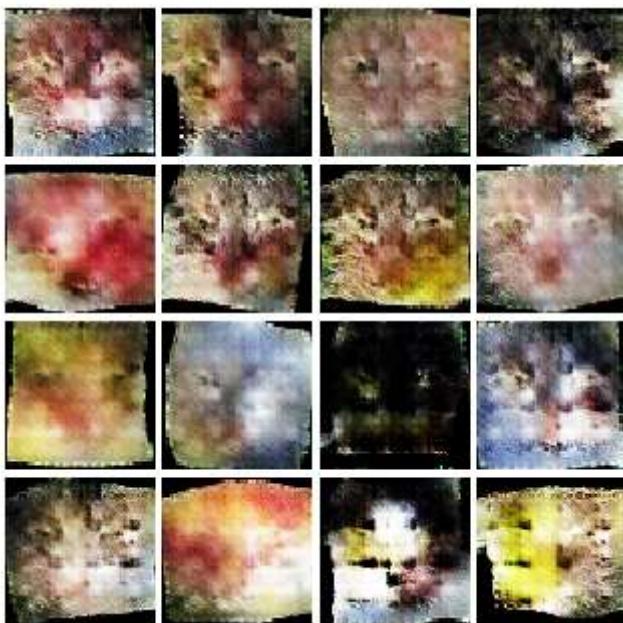
EPOCH: 18

Iter: 8500, D: 0.1721, G: 0.3255



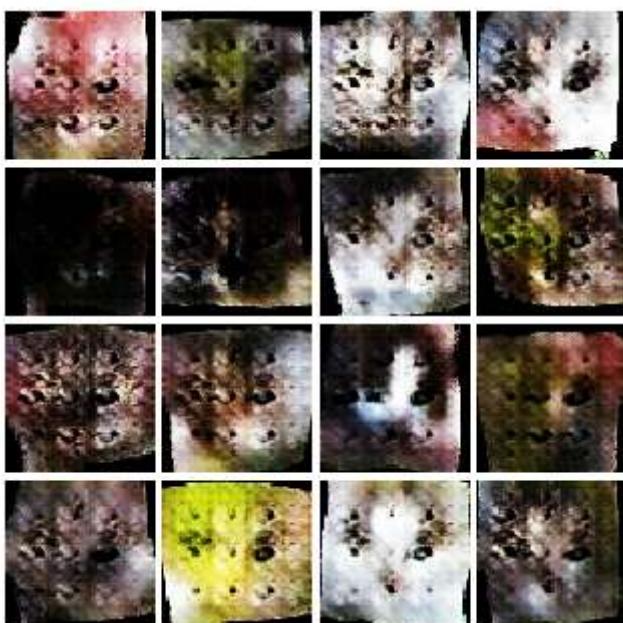


Iter: 8750, D: 0.5474, G:0.2671

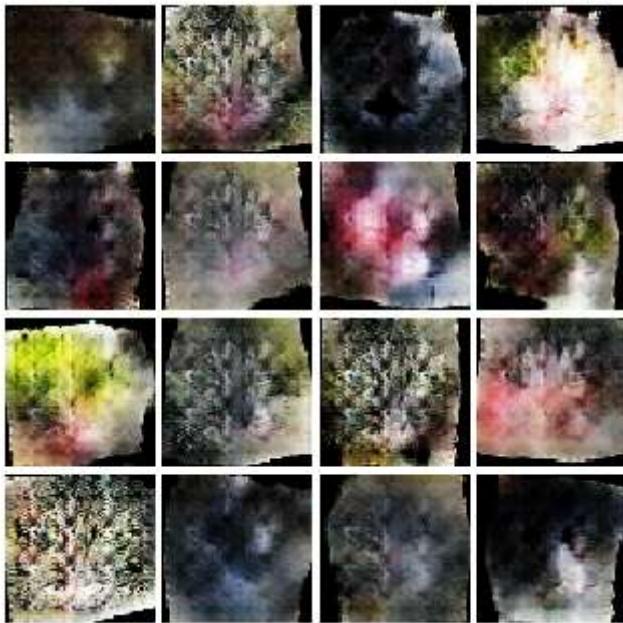


EPOCH: 19

Iter: 9000, D: 0.1768, G:0.268



Iter: 9250, D: 0.2314, G:0.2239



EPOCH: 20

Iter: 9500, D: 0.3198, G:0.2107



Iter: 9750, D: 0.2722, G:0.1819





EPOCH: 21

Iter: 10000, D: 0.3473, G: 0.2727



Iter: 10250, D: 0.2886, G: 0.2131



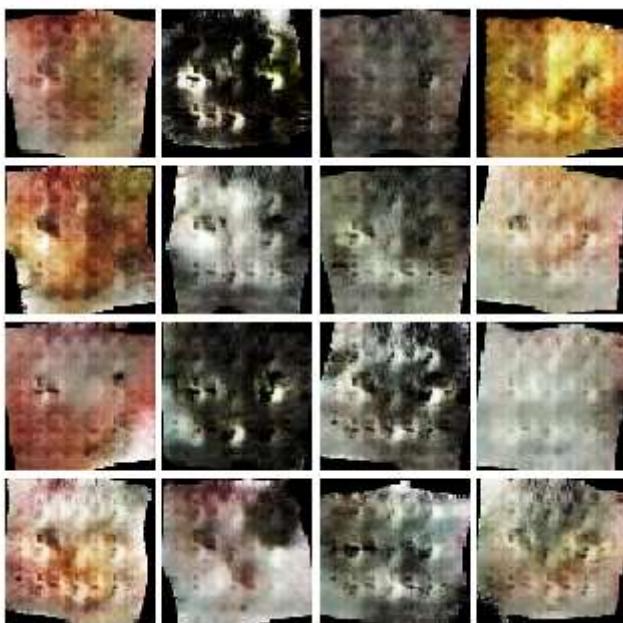
EPOCH: 22

Iter: 10500, D: 0.2288, G: 0.2429





Iter: 10750, D: 0.2085, G: 0.3348



EPOCH: 23

Iter: 11000, D: 0.1942, G: 0.1443



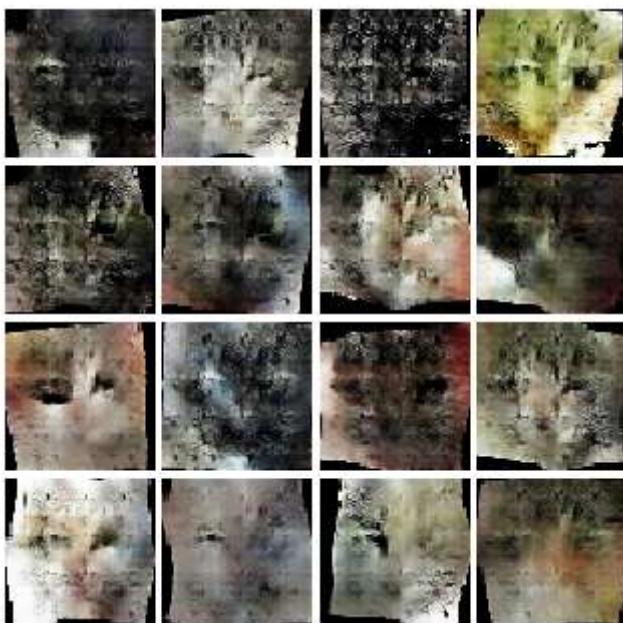


Iter: 11250, D: 0.2767, G:0.2884



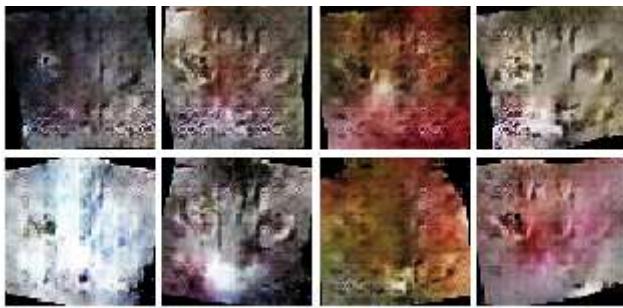
EPOCH: 24

Iter: 11500, D: 0.1252, G:0.3581



Iter: 11750, D: 0.1488, G:0.2507





EPOCH: 25

Iter: 12000, D: 0.2228, G: 0.3441



Iter: 12250, D: 0.1896, G: 0.4519



EPOCH: 26

Iter: 12500, D: 0.2261, G: 0.7038





Iter: 12750, D: 0.2473, G: 0.8095



EPOCH: 27

Iter: 13000, D: 0.1631, G: 0.5141





Iter: 13250, D: 0.2342, G: 0.4416



EPOCH: 28

Iter: 13500, D: 0.211, G: 0.3536



Iter: 13750, D: 0.4228, G: 1.301





EPOCH: 29

Iter: 14000, D: 0.31, G: 0.2028



Iter: 14250, D: 0.1277, G: 0.4235



EPOCH: 30

Iter: 14500, D: 0.26, G: 0.799



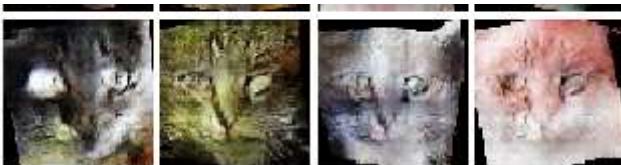
Iter: 14750, D: 0.2193, G: 0.6028



EPOCH: 31

Iter: 15000, D: 0.1245, G: 0.5017





Iter: 15250, D: 0.2277, G: 0.2731



EPOCH: 32

Iter: 15500, D: 0.1135, G: 0.3613



EPOCH: 33

Iter: 15750, D: 0.07075, G: 0.3894





Iter: 16000, D: 0.1484, G: 0.6708



EPOCH: 34

Iter: 16250, D: 0.2082, G: 0.133



Iter: 16500, D: 0.3036, G:0.2128



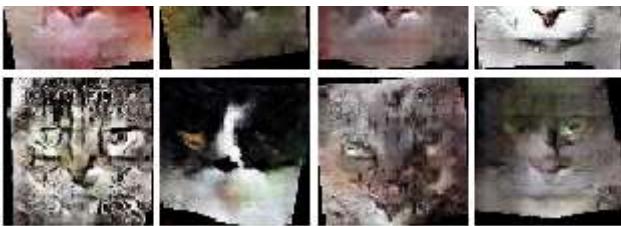
EPOCH: 35

Iter: 16750, D: 0.2497, G:0.3855



Iter: 17000, D: 0.176, G:0.4333





EPOCH: 36

Iter: 17250, D: 0.03817, G:0.4479



Iter: 17500, D: 0.07675, G:0.3079



EPOCH: 37

Iter: 17750, D: 0.08682, G:0.4113





Iter: 18000, D: 0.1552, G: 0.2979



EPOCH: 38

Iter: 18250, D: 0.0697, G: 0.4067



Iter: 18500, D: 0.1042, G:0.5755



EPOCH: 39

Iter: 18750, D: 0.09618, G:0.4593



Iter: 19000, D: 0.06422, G:0.4242





EPOCH: 40

Iter: 19250, D: 0.03282, G:0.4334



Iter: 19500, D: 0.04452, G:0.7536



EPOCH: 41

Iter: 19750, D: 0.06267, G:0.636





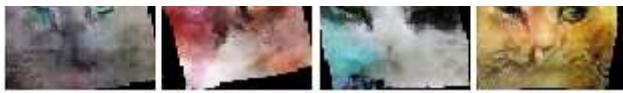
Iter: 20000, D: 0.1207, G: 0.6796



EPOCH: 42

Iter: 20250, D: 0.04635, G: 0.4948





Iter: 20500, D: 0.08243, G:0.4496



EPOCH: 43

Iter: 20750, D: 0.1925, G:0.9108



Iter: 21000, D: 0.1914, G:0.3235





EPOCH: 44

Iter: 21250, D: 0.09725, G:0.4987



Iter: 21500, D: 0.04754, G:0.5303



EPOCH: 45

Iter: 21750, D: 0.02453, G:0.3757



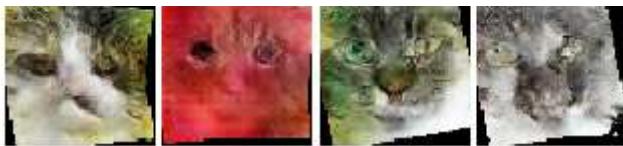
Iter: 22000, D: 0.02569, G:0.6016



EPOCH: 46

Iter: 22250, D: 0.1362, G:0.9052





Iter: 22500, D: 0.07824, G:0.4224



EPOCH: 47

Iter: 22750, D: 0.09832, G:0.3758



Iter: 23000, D: 0.08666, G:0.2883





EPOCH: 48

Iter: 23250, D: 0.09723, G:0.4557



Iter: 23500, D: 0.04062, G:0.7224



EPOCH: 49

Iter: 23750, D: 0.06536, G:0.4853



Iter: 24000, D: 0.03182, G:0.4533



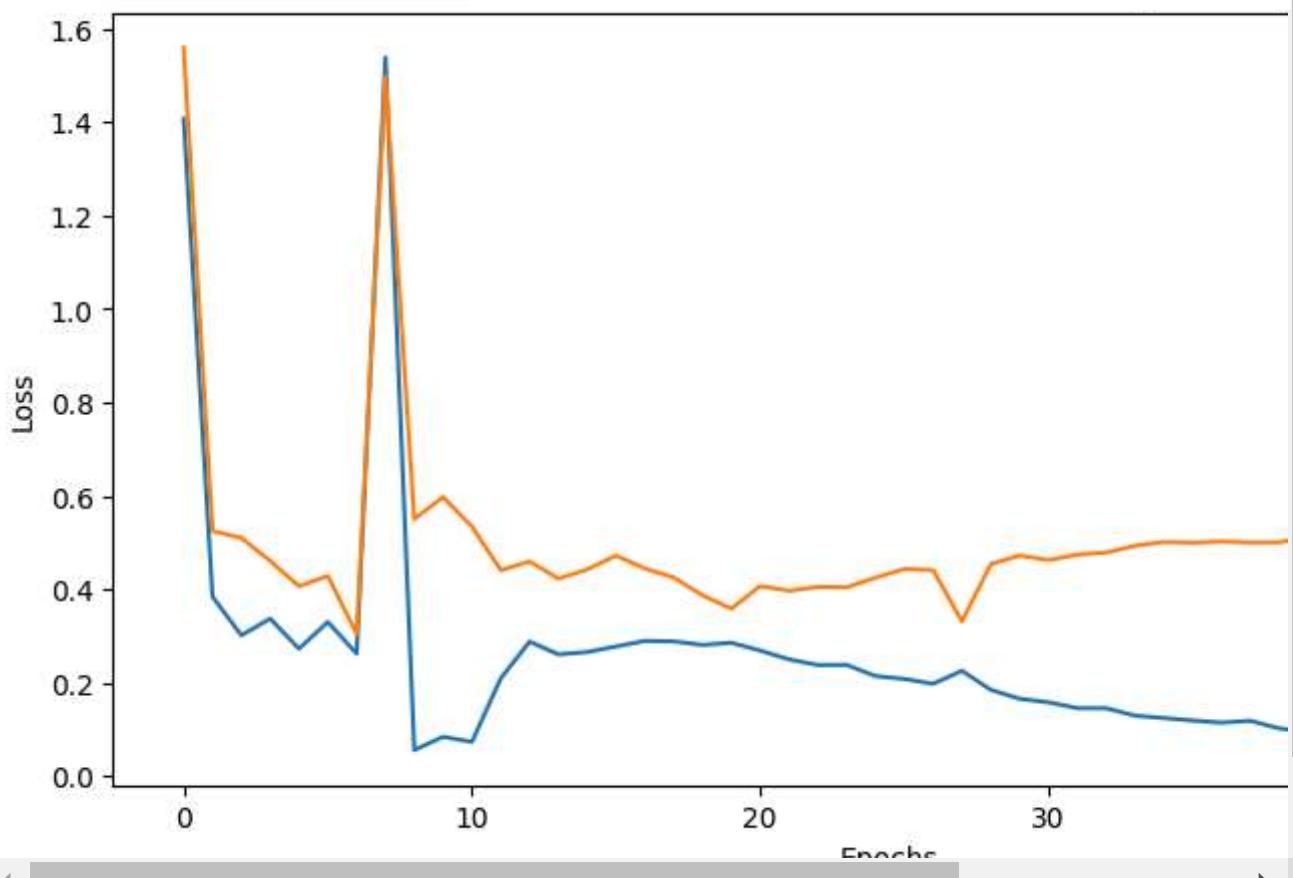
EPOCH: 50

Iter: 24250, D: 0.07781, G:0.3245





Iter: 24500, D: 0.07277, G: 0.3537



▼ WGAN Loss

```
from losses import wg_generator_loss, wg_discriminator_loss

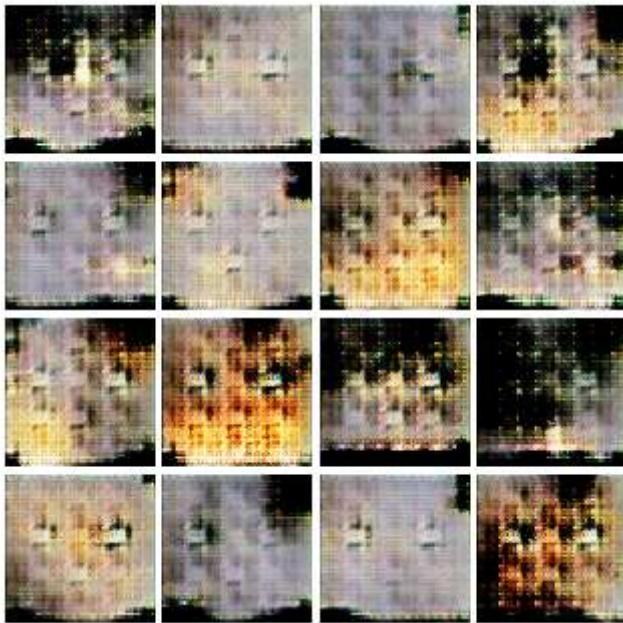
D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)

D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

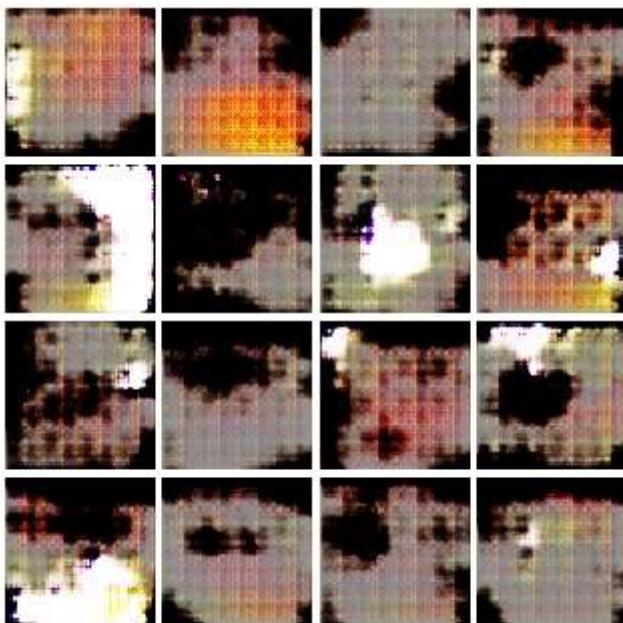
# ls-gan
train(D, G, D_optimizer, G_optimizer, wg_discriminator_loss,
      wg_generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

EPOCH: 1

Iter: 0, D: -1.06, G:6.693

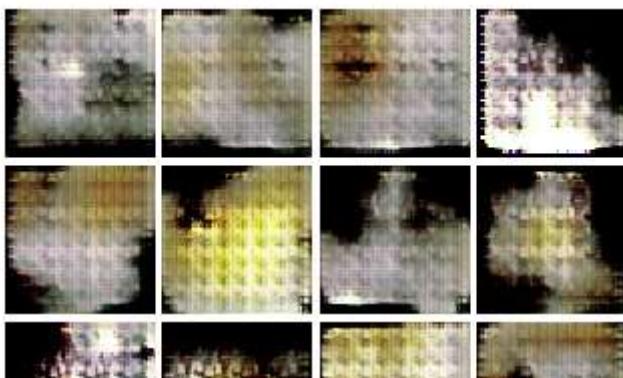


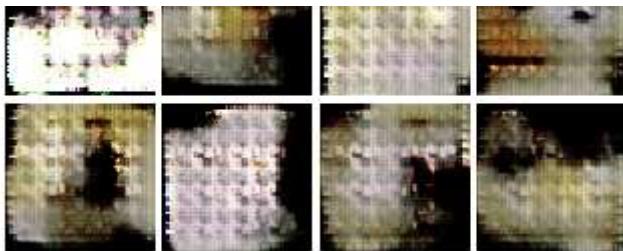
Iter: 250, D: -789.3, G:3.611e+03



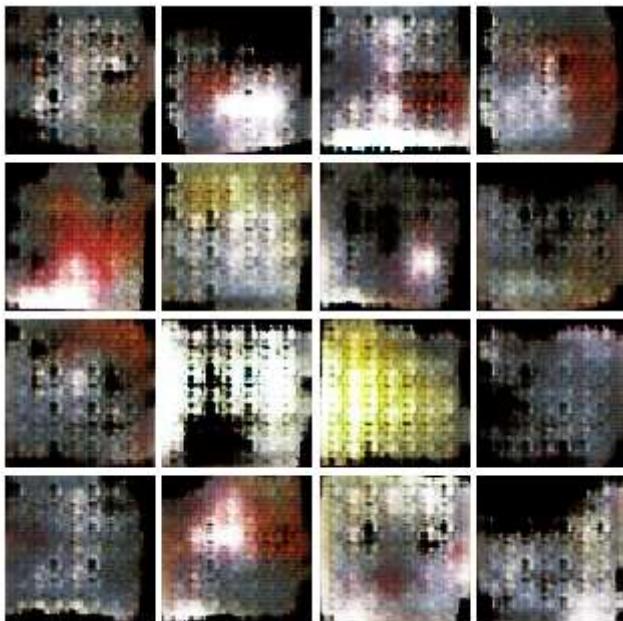
EPOCH: 2

Iter: 500, D: -1.763e+04, G:-1.548e+03



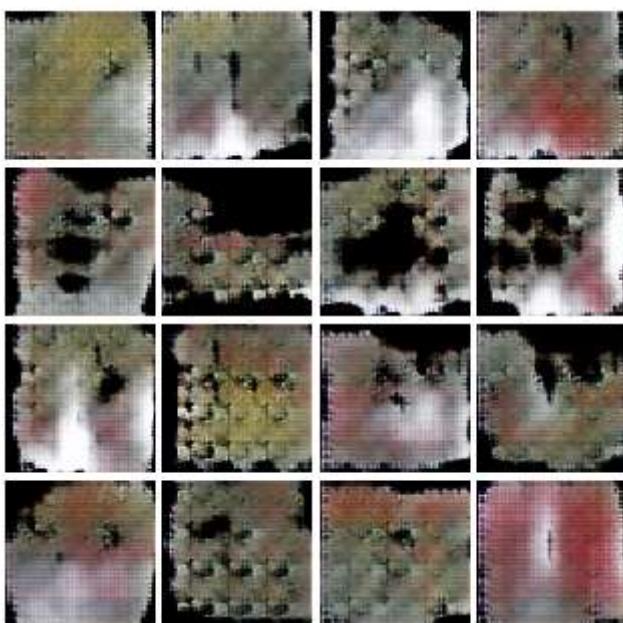


Iter: 750, D: -7.65e+03, G:2.425e+04



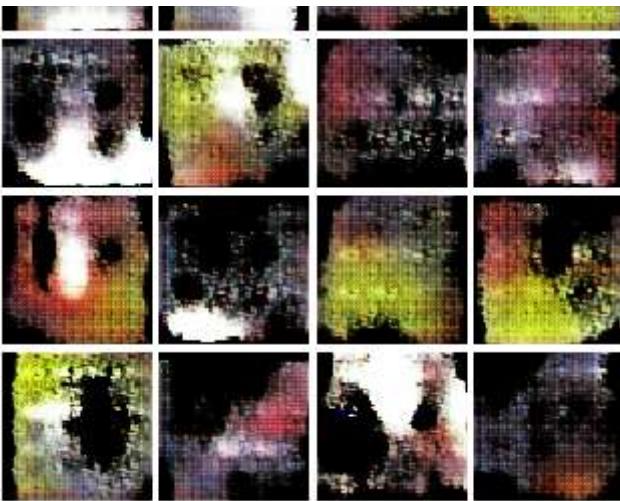
EPOCH: 3

Iter: 1000, D: -8.674e+04, G:4.58e+04



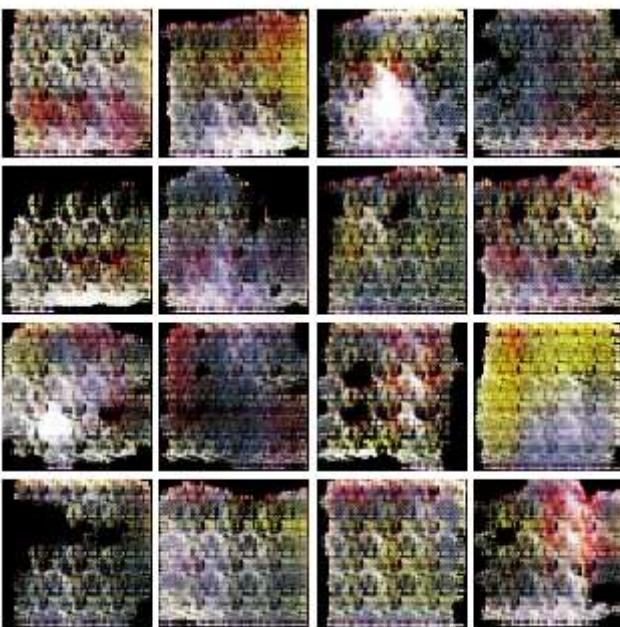
Iter: 1250, D: -1.3e+05, G:6.787e+04



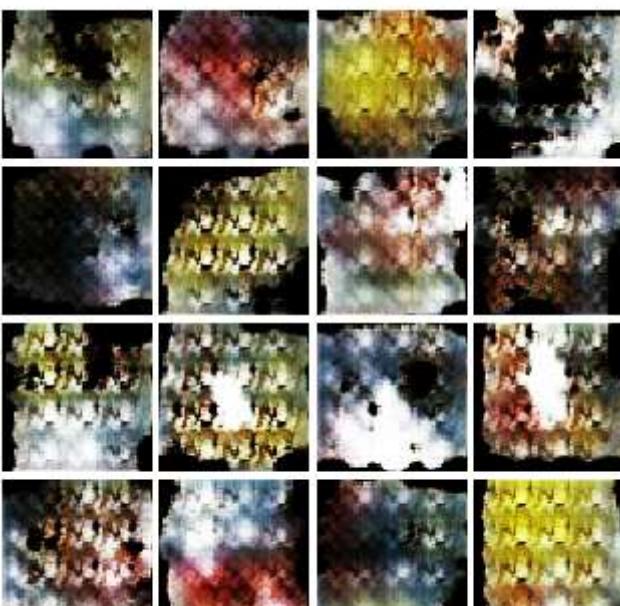


EPOCH: 4

Iter: 1500, D: -1.467e+04, G:-1.511e+04

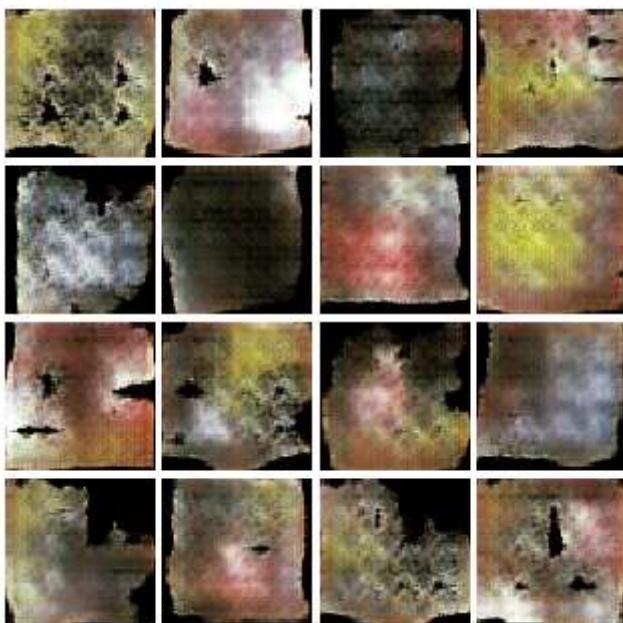
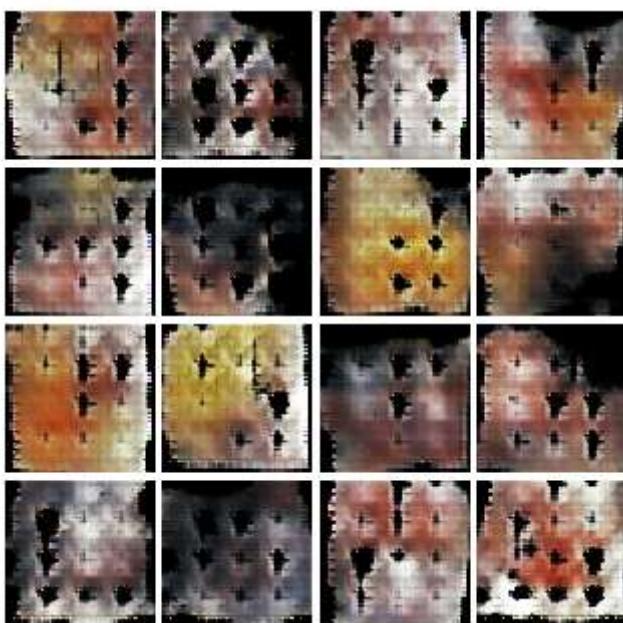


Iter: 1750, D: -2.216e+05, G:1.173e+05



EPOCH: 5

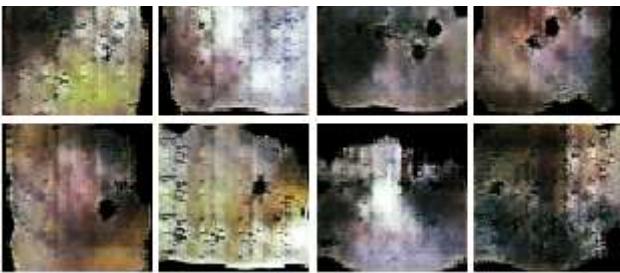
Iter: 2000, D: -2.967e+05, G:1.557e+05



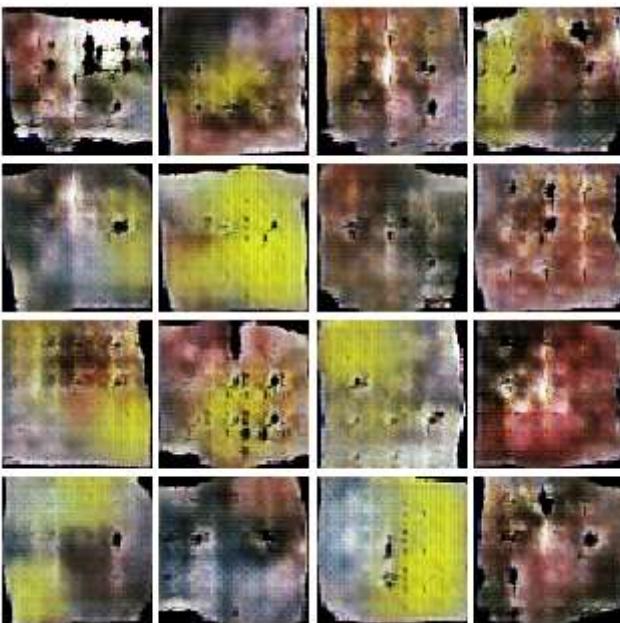
EPOCH: 6

Iter: 2500, D: -4.003e+05, G:2.146e+05





Iter: 2750, D: -4.452e+05, G:2.366e+05



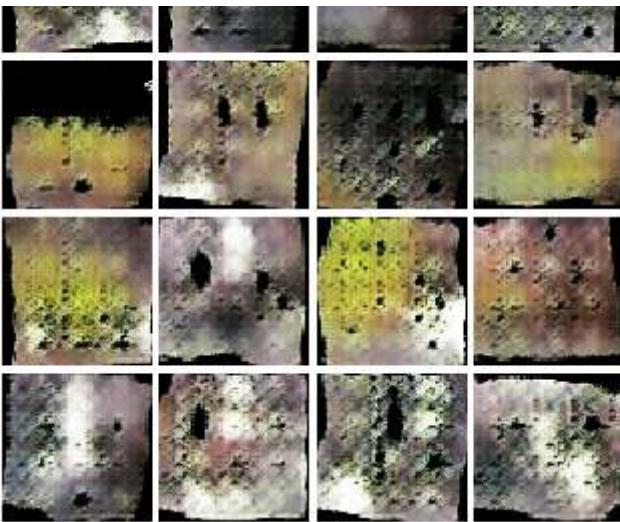
EPOCH: 7

Iter: 3000, D: -5.89e+05, G:3.091e+05



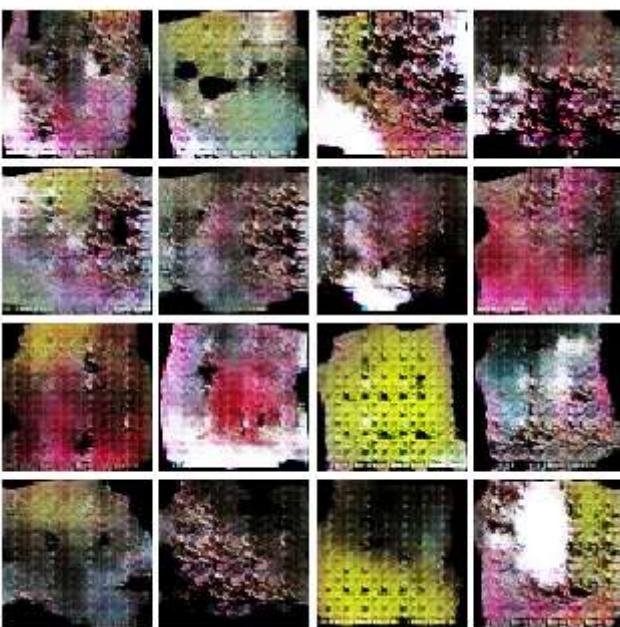
Iter: 3250, D: -6.439e+05, G:3.462e+05





EPOCH: 8

Iter: 3500, D: -7.846e+05, G:4.099e+05



Iter: 3750, D: -8.375e+04, G:3.706e+05



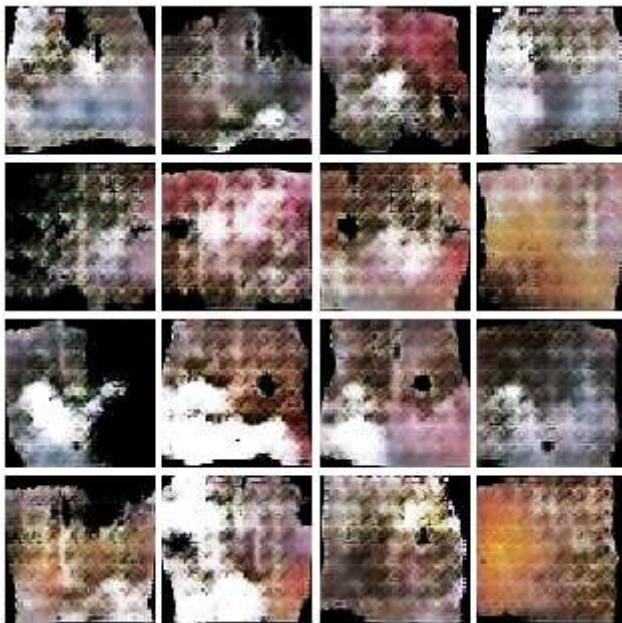


EPOCH: 9

Iter: 4000, D: -9.725e+05, G:5.095e+05

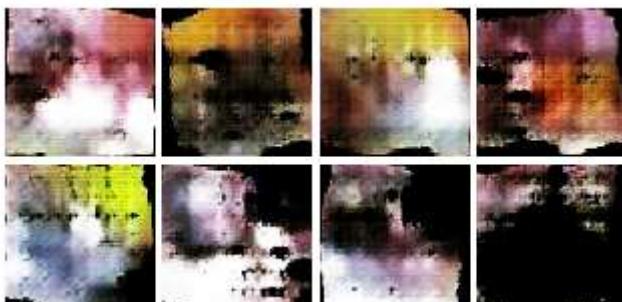


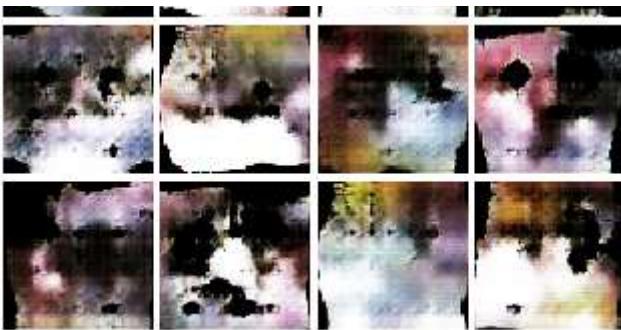
Iter: 4250, D: -1.067e+06, G:5.636e+05



EPOCH: 10

Iter: 4500, D: -2.109e+05, G:4.849e+05





Iter: 4750, D: -1.302e+06, G:6.811e+05



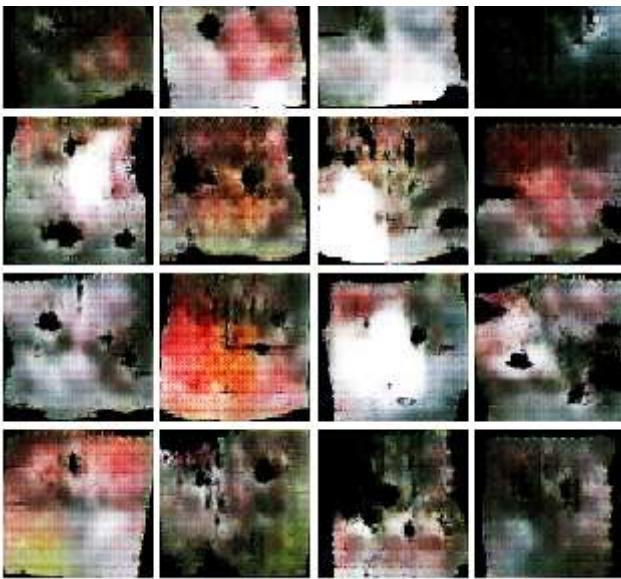
EPOCH: 11

Iter: 5000, D: -1.448e+06, G:7.581e+05



Iter: 5250, D: -1.938e+05, G:1.64e+05





EPOCH: 12

Iter: 5500, D: -1.25e+06, G:7.536e+05



Iter: 5750, D: -6.318e+05, G:8.336e+05





EPOCH: 13

Iter: 6000, D: -1.79e+06, G:9.638e+05



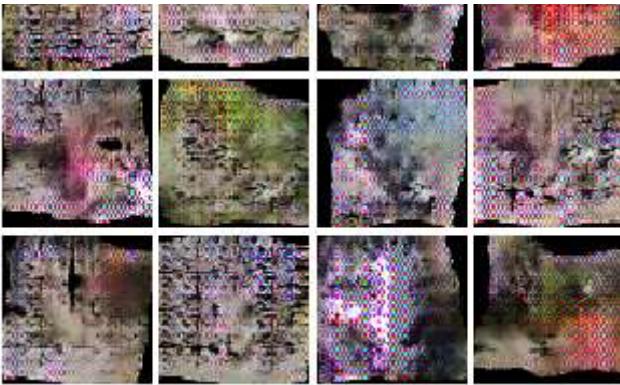
Iter: 6250, D: -2.076e+06, G:1.087e+06



EPOCH: 14

Iter: 6500, D: -2.247e+06, G:1.175e+06



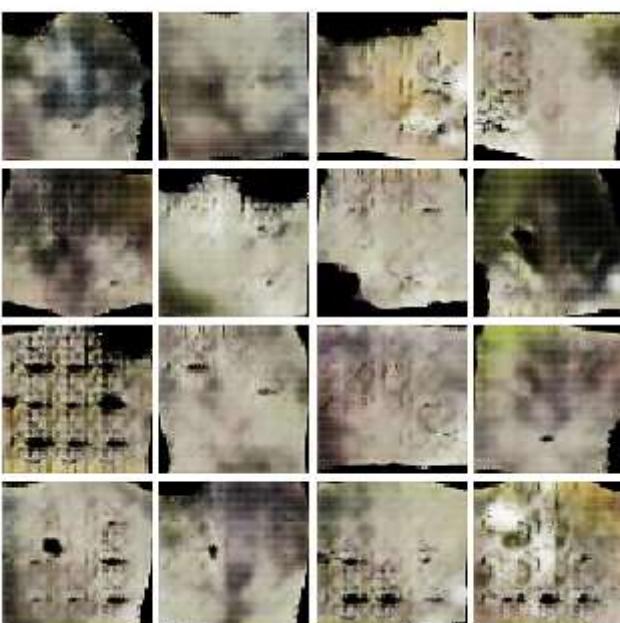


Iter: 6750, D: -2.416e+06, G:1.264e+06

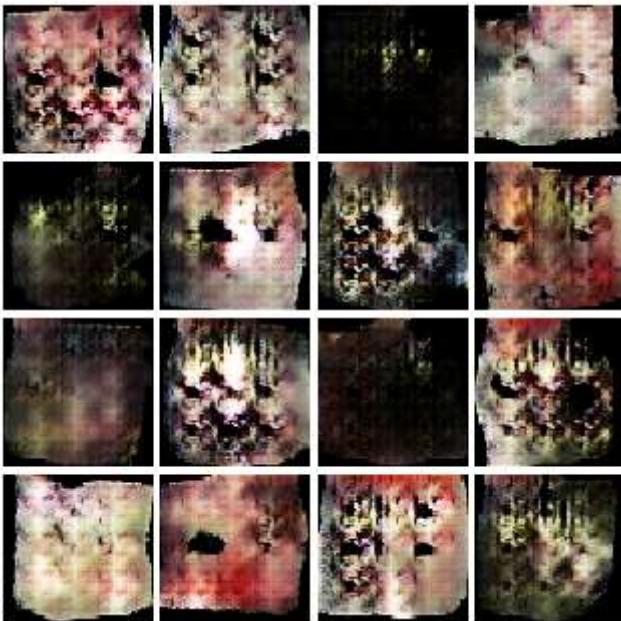


EPOCH: 15

Iter: 7000, D: -2.588e+06, G:1.356e+06



Iter: 7250, D: -1.511e+05, G:-6.935e+05



EPOCH: 16

Iter: 7500, D: -2.408e+06, G:1.416e+06



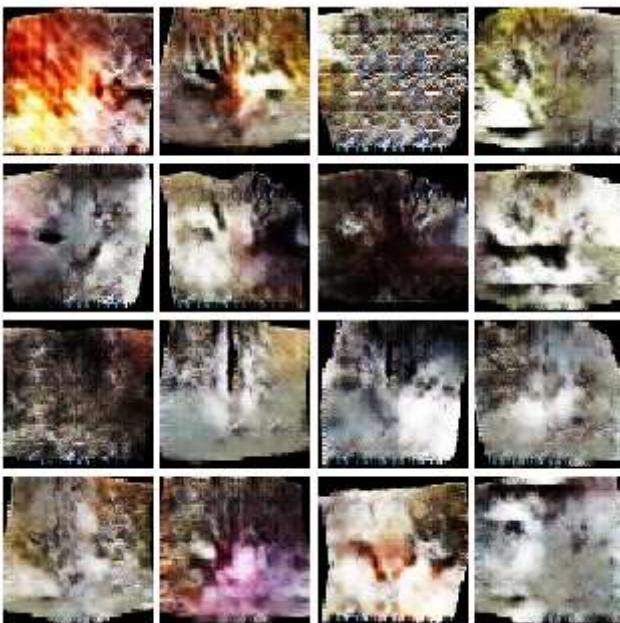
Iter: 7750, D: -2.776e+06, G:1.322e+06



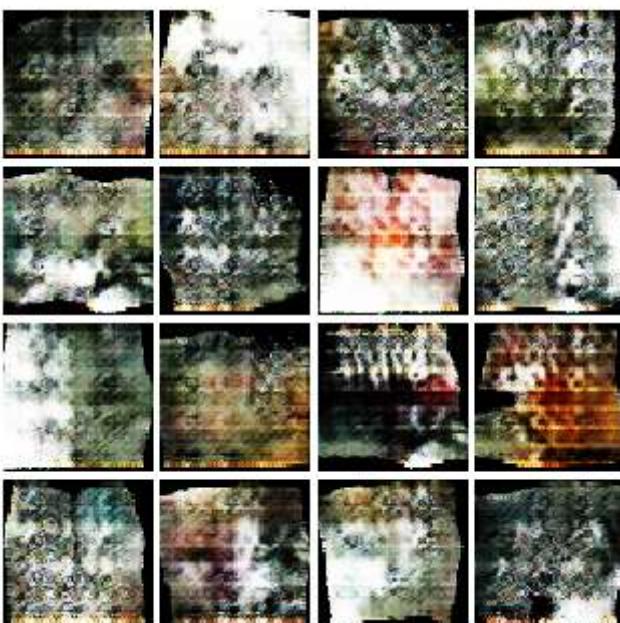


EPOCH: 17

Iter: 8000, D: -3.151e+06, G:1.654e+06



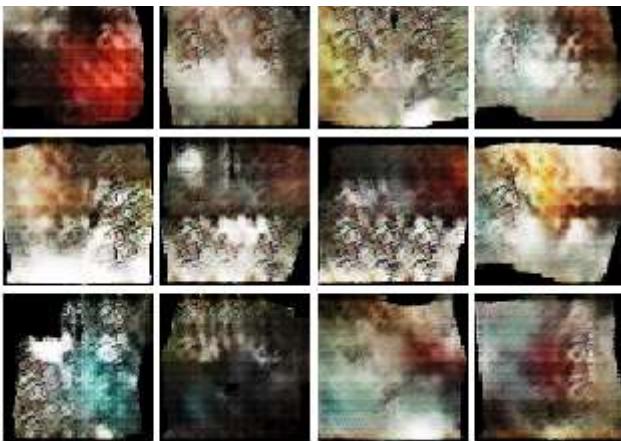
Iter: 8250, D: -3.345e+06, G:1.753e+06



EPOCH: 18

Iter: 8500, D: -3.528e+06, G:1.861e+06



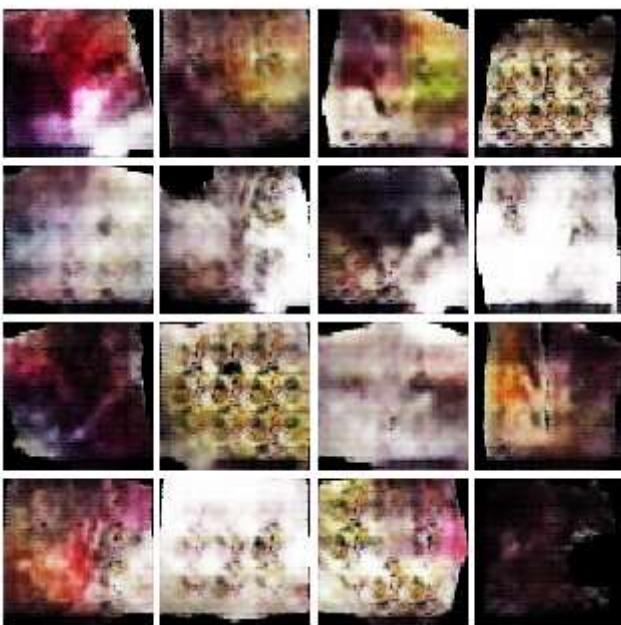


Iter: 8750, D: -3.756e+06, G:1.971e+06



EPOCH: 19

Iter: 9000, D: -3.941e+06, G:2.061e+06



Iter: 9250, D: -4.173e+06, G:2.19e+06



EPOCH: 20

Iter: 9500, D: -4.382e+06, G:2.3e+06



Iter: 9750, D: -4.591e+06, G:2.409e+06





EPOCH: 21

Iter: 10000, D: -4.593e+06, G: 2.434e+06



Iter: 10250, D: -4.895e+06, G: 2.582e+06



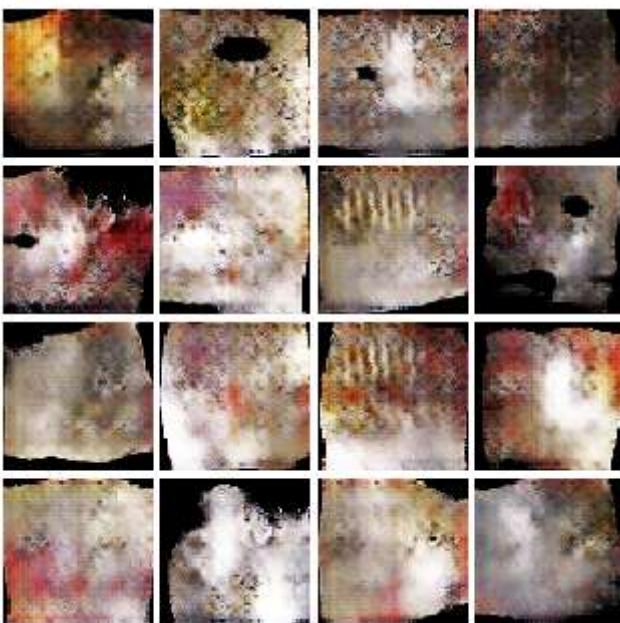
EPOCH: 22

Iter: 10500, D: -5.158e+06, G: 2.711e+06



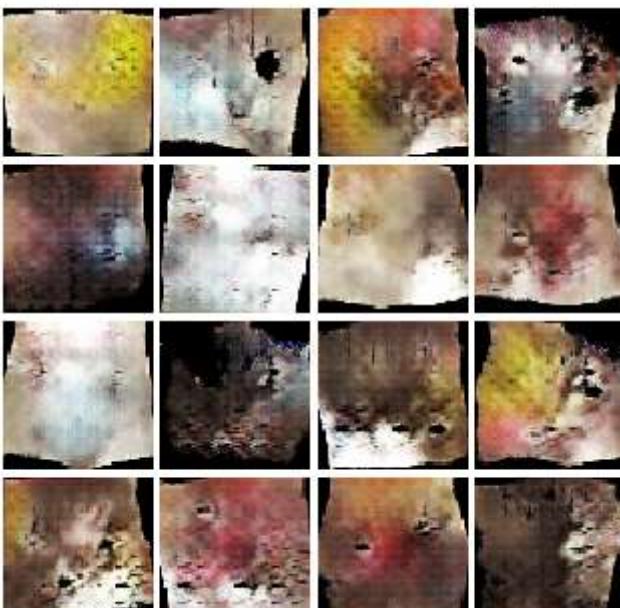


Iter: 10750, D: -5.311e+06, G:2.803e+06

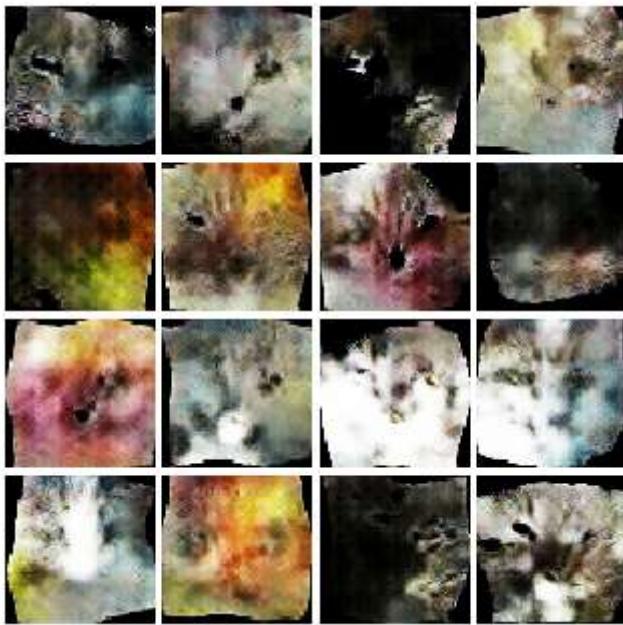


EPOCH: 23

Iter: 11000, D: -1.321e+06, G:2.152e+06



Iter: 11250, D: -4.984e+06, G:2.826e+06



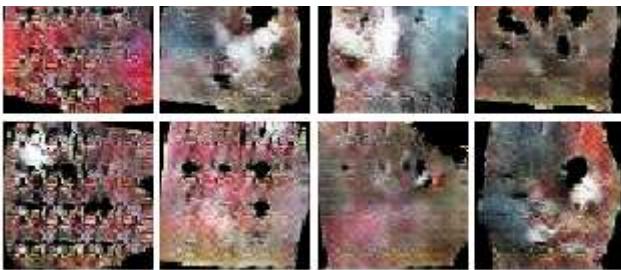
EPOCH: 24

Iter: 11500, D: -5.916e+06, G:3.113e+06



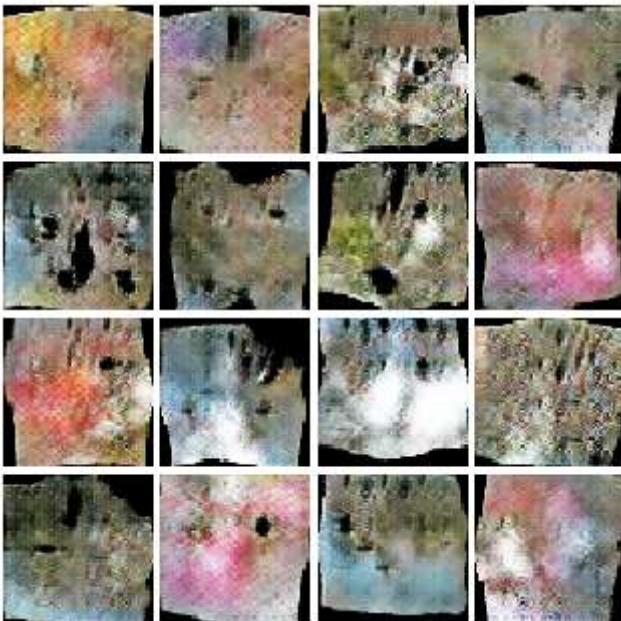
Iter: 11750, D: -6.187e+06, G:3.256e+06



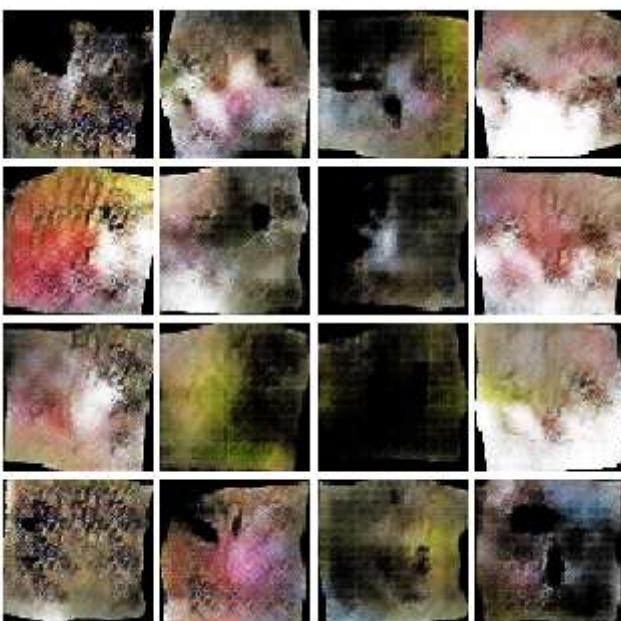


EPOCH: 25

Iter: 12000, D: -6.405e+06, G:3.382e+06



Iter: 12250, D: -6.733e+06, G:3.545e+06



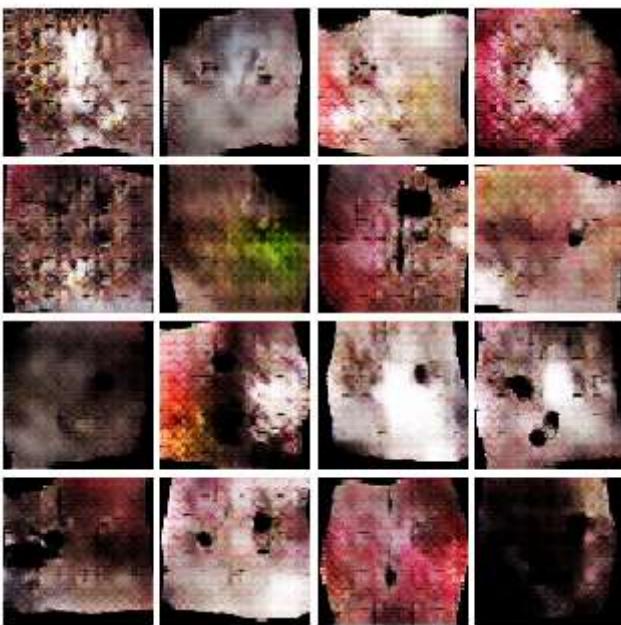
EPOCH: 26

Iter: 12500, D: -6.966e+06, G:3.681e+06





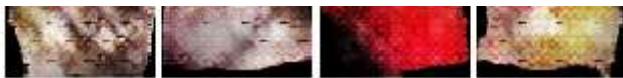
Iter: 12750, D: -7.249e+06, G:3.815e+06



EPOCH: 27

Iter: 13000, D: -7.521e+06, G:3.957e+06



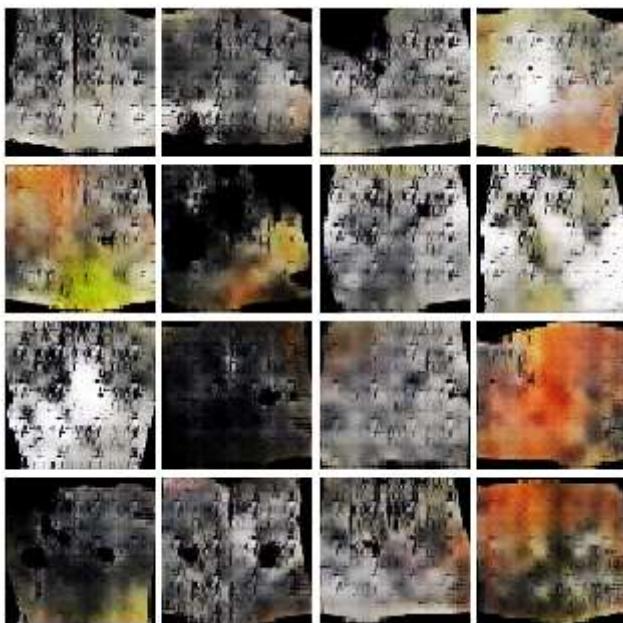


Iter: 13250, D: -7.767e+06, G:4.087e+06



EPOCH: 28

Iter: 13500, D: -8.01e+06, G:4.225e+06



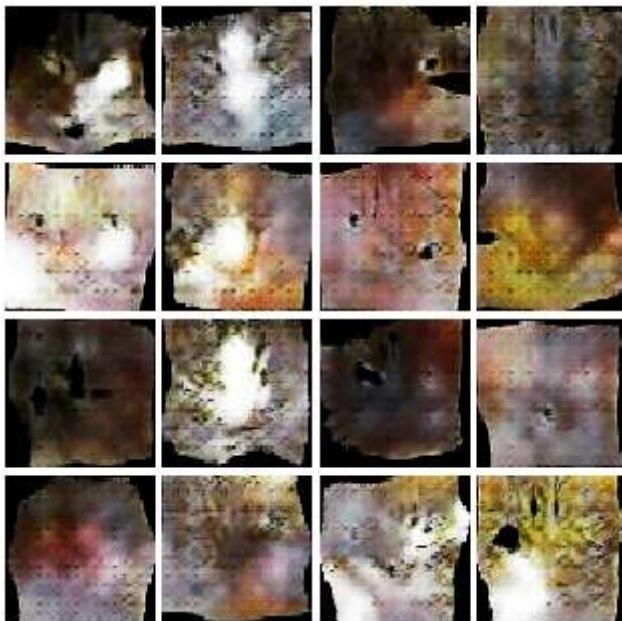
Iter: 13750, D: -2.116e+06, G:2.05e+06





EPOCH: 29

Iter: 14000, D: -7.864e+06, G:4.285e+06



Iter: 14250, D: -8.628e+06, G:4.562e+06



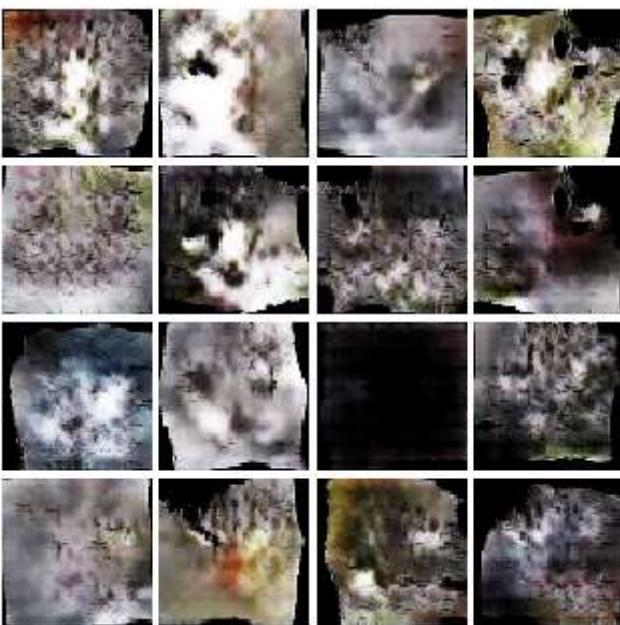
EPOCH: 30

Iter: 14500, D: -8.967e+06, G:4.719e+06





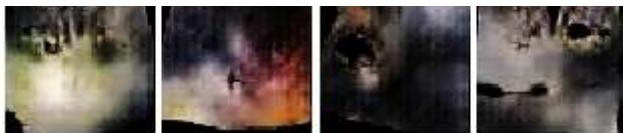
Iter: 14750, D: -4.127e+06, G:3.559e+06



EPOCH: 31

Iter: 15000, D: -3.577e+06, G:-1.131e+06





Iter: 15250, D: -8.457e+06, G: 4.861e+06



EPOCH: 32

Iter: 15500, D: -9.983e+06, G: 5.264e+06



EPOCH: 33

Iter: 15750, D: -1.032e+07, G: 5.446e+06



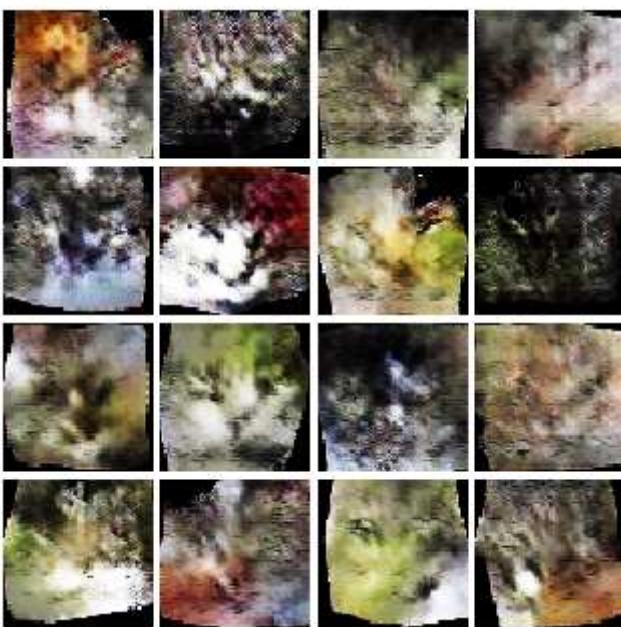


Iter: 16000, D: -7.395e+06, G:4.499e+06



EPOCH: 34

Iter: 16250, D: -9.398e+06, G:4.252e+06



File 16500 16700 16750 16780 16800

Iter: 16500, D: -1.116e+07, G:5.8/8e+06



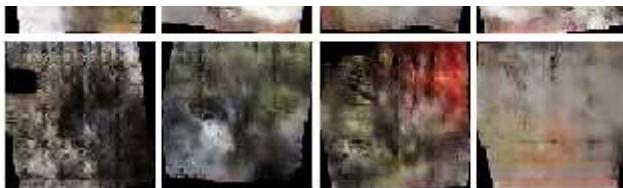
EPOCH: 35

Iter: 16750, D: -5.908e+05, G:5.681e+06



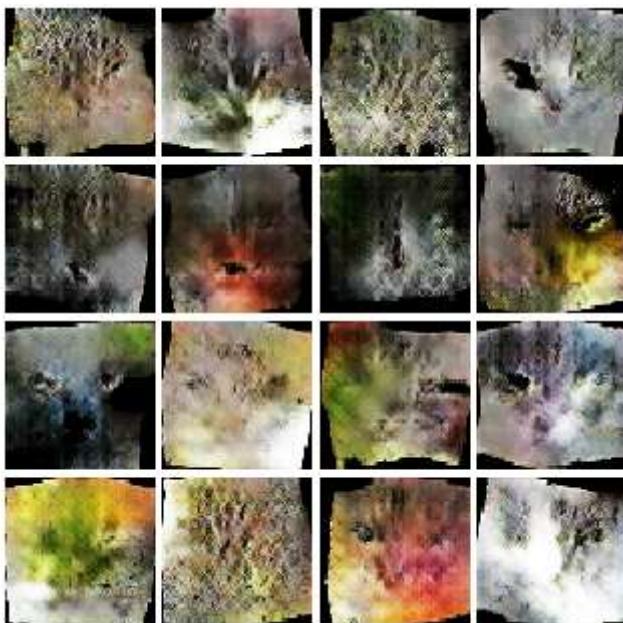
Iter: 17000, D: -1.182e+07, G:6.252e+06





EPOCH: 36

Iter: 17250, D: -1.19e+07, G:6.329e+06



Iter: 17500, D: -1.255e+07, G:6.617e+06



EPOCH: 37

Iter: 17750, D: -1.292e+07, G:6.81e+06





Iter: 18000, D: -1.328e+07, G:7.004e+06

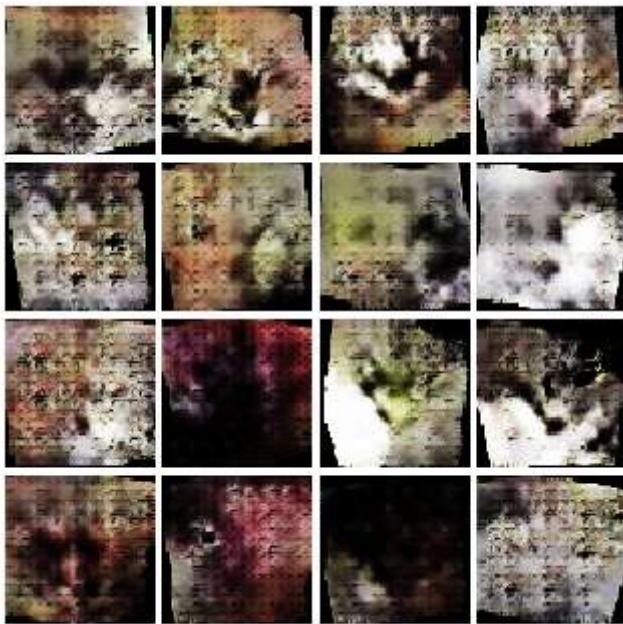


EPOCH: 38

Iter: 18250, D: -1.362e+07, G:7.186e+06



Iter: 18500, D: -1.401e+07, G:7.394e+06



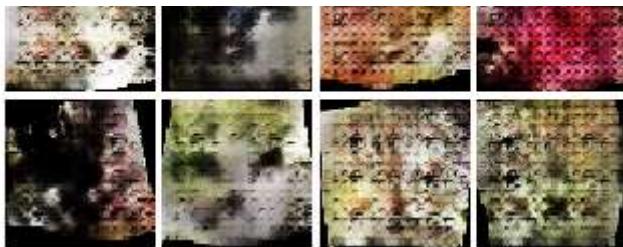
EPOCH: 39

Iter: 18750, D: -1.439e+07, G:7.588e+06



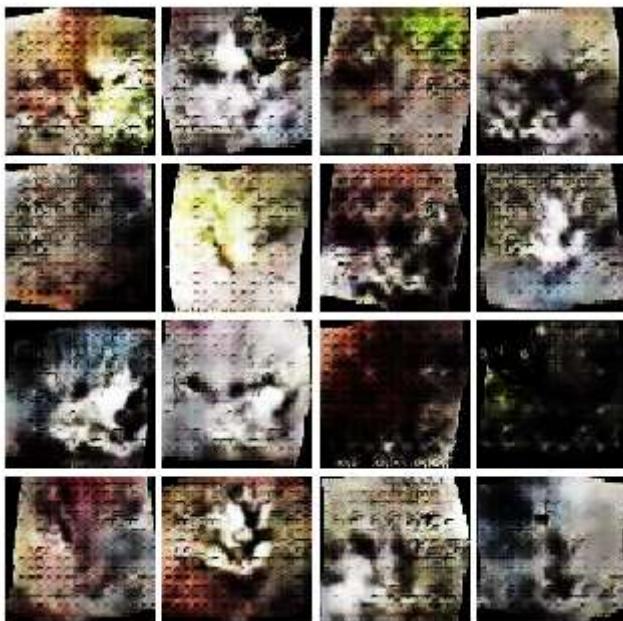
Iter: 19000, D: -1.475e+07, G:7.781e+06



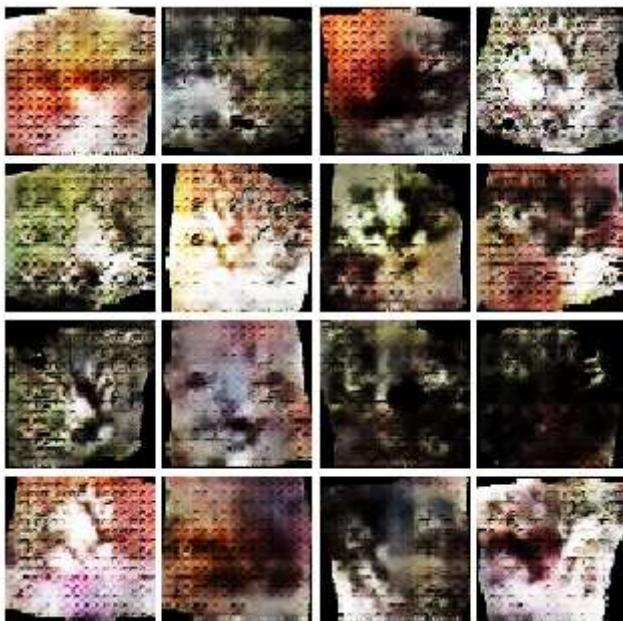


EPOCH: 40

Iter: 19250, D: -1.511e+07, G:7.972e+06



Iter: 19500, D: -1.548e+07, G:8.167e+06



EPOCH: 41

Iter: 19750, D: -1.585e+07, G:8.362e+06





Iter: 20000, D: -1.619e+07, G:8.515e+06



EPOCH: 42

Iter: 20250, D: -1.661e+07, G:8.767e+06



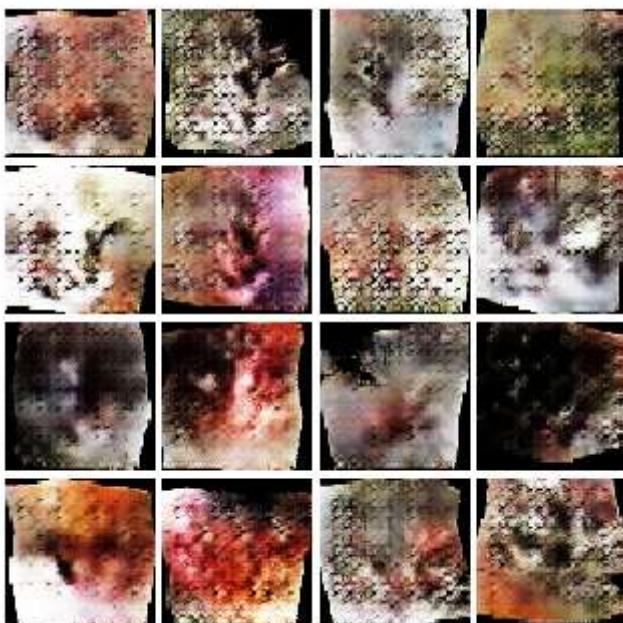


Iter: 20500, D: -1.7e+07, G:8.969e+06

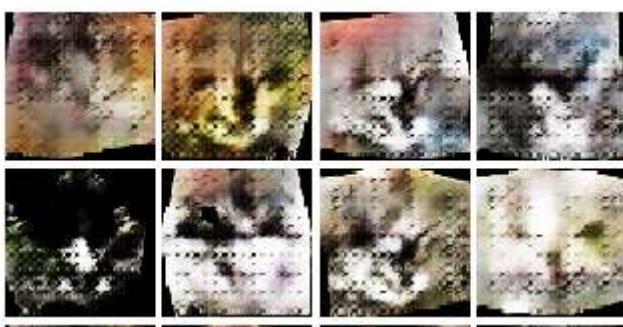


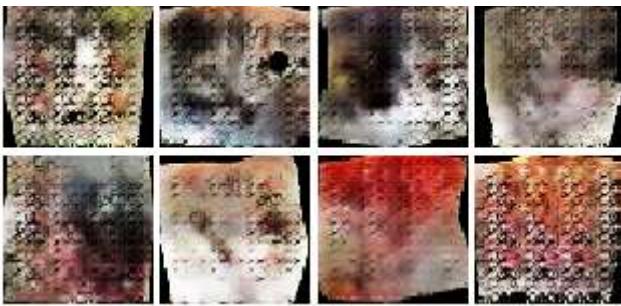
EPOCH: 43

Iter: 20750, D: -1.738e+07, G:9.174e+06



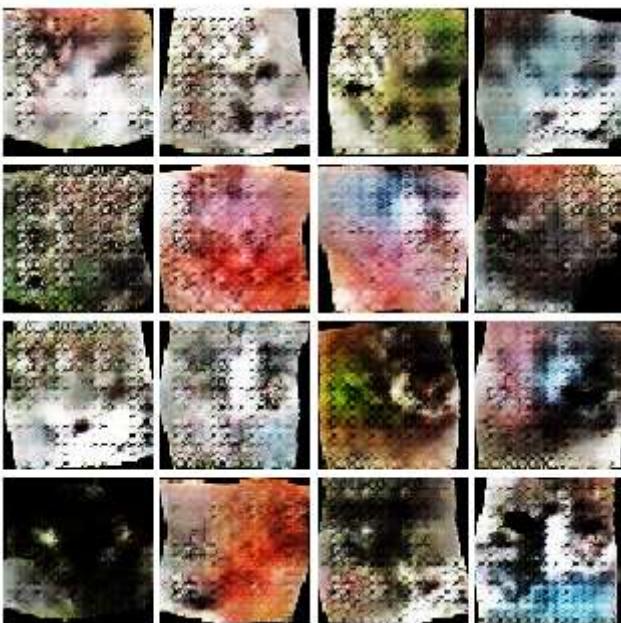
Iter: 21000, D: -1.777e+07, G:9.381e+06



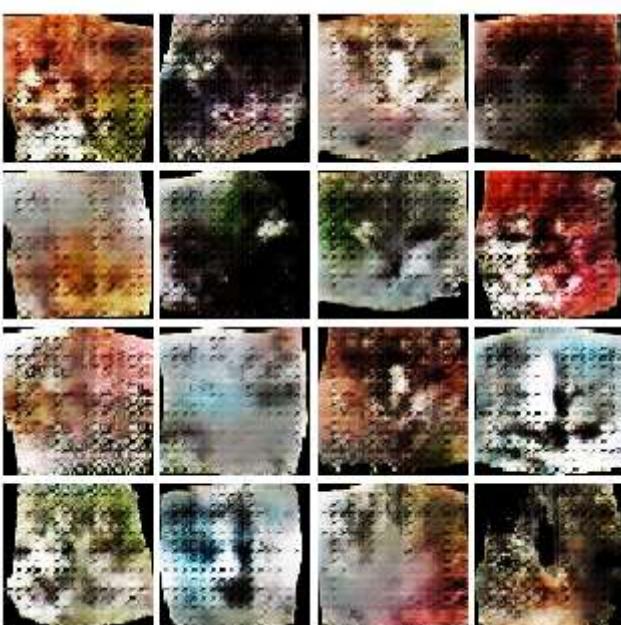


EPOCH: 44

Iter: 21250, D: -1.816e+07, G: 9.589e+06



Iter: 21500, D: -1.856e+07, G: 9.799e+06



EPOCH: 45

Iter: 21750, D: -1.897e+07, G: 1.001e+07





Iter: 22000, D: -1.937e+07, G:1.023e+07



EPOCH: 46

Iter: 22250, D: -1.978e+07, G:1.044e+07



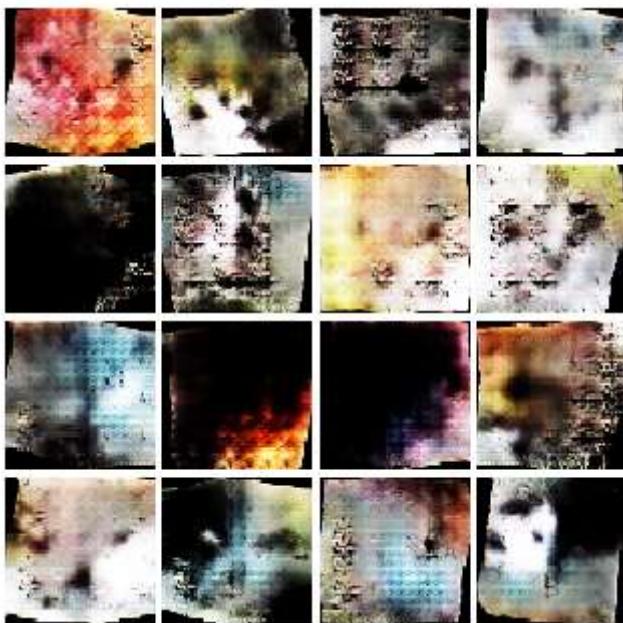


Iter: 22500, D: -2.019e+07, G:1.066e+07



EPOCH: 47

Iter: 22750, D: -2.06e+07, G:1.088e+07



Iter: 23000, D: -2.101e+07, G:1.11e+07





EPOCH: 48

Iter: 23250, D: -2.145e+07, G:1.133e+07

